

# Experiences in Delivering Spark as a Service

**July 27 2016**

**Michael Feiman, Khalid Ahmed  
STSM, IBM Spectrum Computing**

# Agenda

1

Bluemix Spark Cloud Technical Challenges

2

Bluemix Spark Cloud Architecture & Implementation

3

IBM – Spark & Mesos Community

# IBM Spectrum Computing

## *Infrastructure software for high performance applications*

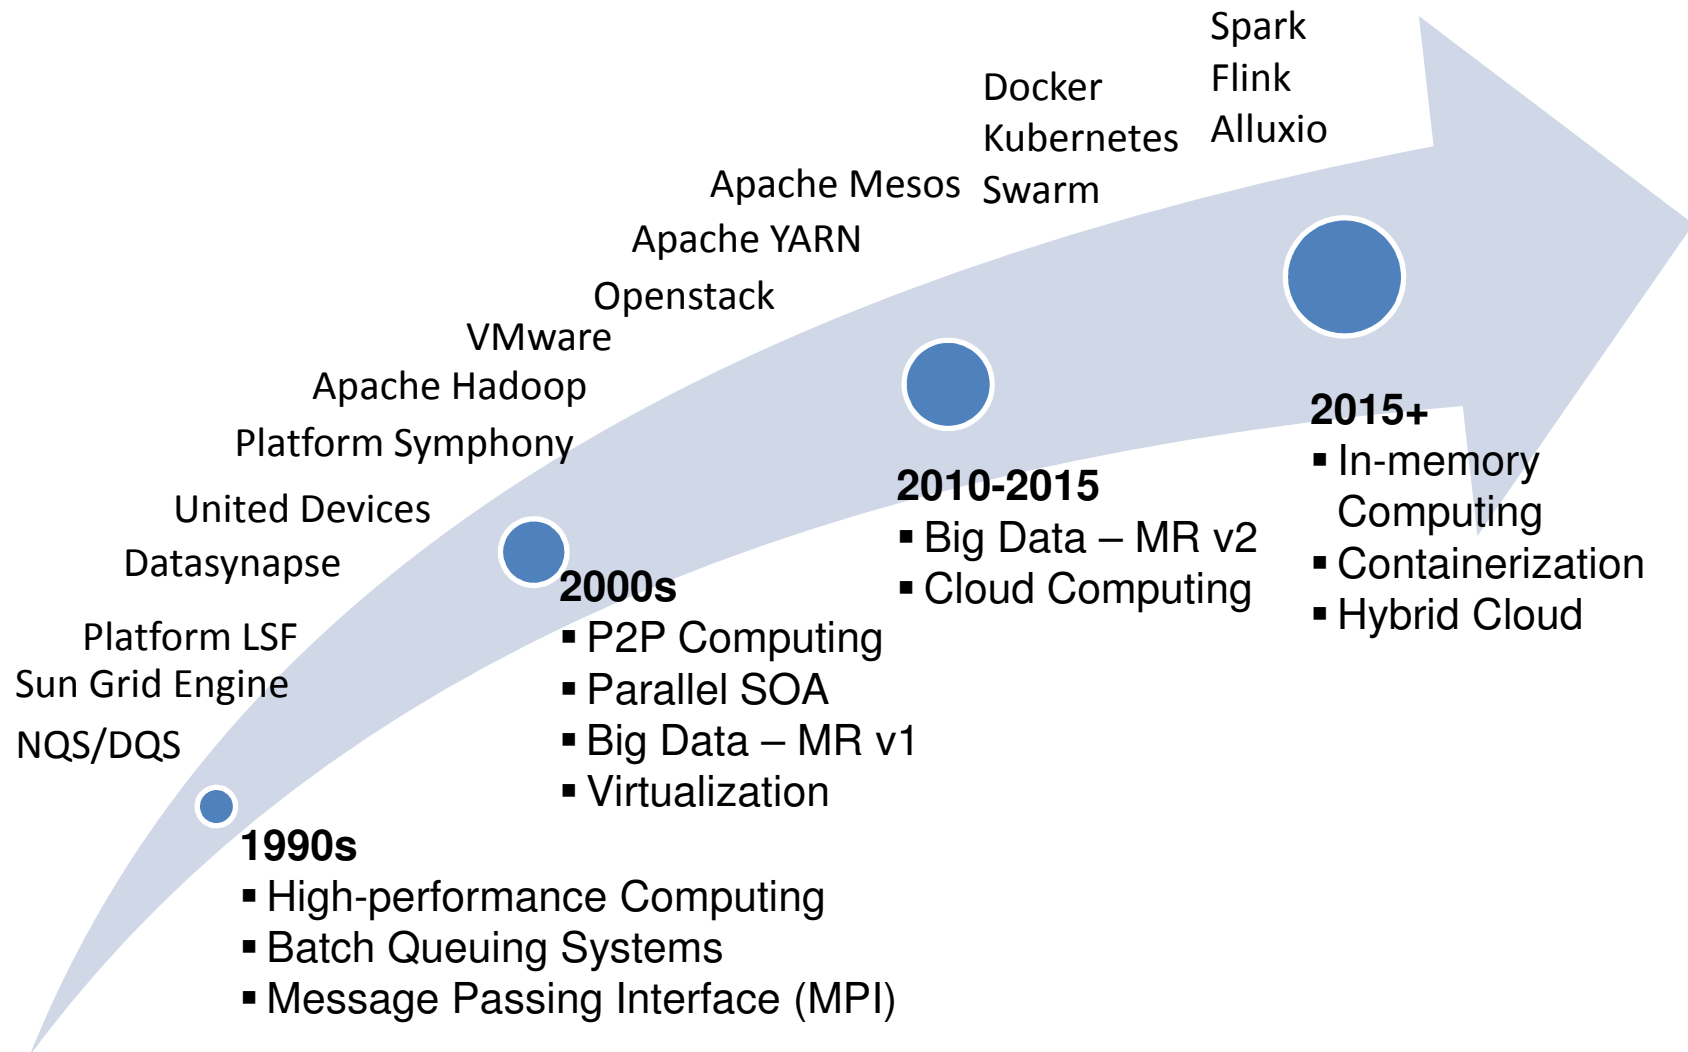
- Acquired by IBM in 2012
- 20 years managing distributed scale-out systems with 2000+ customers in many industries
- Market leading workload, resource and cluster management
- Unmatched scalability (small clusters to global grids) and enterprise production-proven reliability
- Heterogeneous environments – x86 and Power plus 3rd party systems, virtual and bare metal, accelerators / GPU, cloud, etc.
- Shared services for both compute and data intensive workloads

23 of 30 largest  
commercial  
enterprises

Over 5M CPUs under  
management

60% of top  
financial services  
companies

# History of Distributed Cluster Management



# Agenda

1

Bluemix Spark Cloud Technical  
Challenges

2

Bluemix Spark Cloud Architecture &  
Implementation

3

IBM – Spark & OSS Community

# IBM's vision for IBM Analytics for Apache Spark (Spark-as-a-Service)

We make Spark  
**ACCESSIBLE** and **USEFUL**

# Spark is Important!

## Standalone

Spark as a Service

IBM Spectrum  
Conductor with  
Spark (on-prem)

## Within Platforms

IBM Open Platform  
(w/ Spark)

BigInsights (w/  
Spark)

IBM Streams

...many others  
underway

## Within Solutions

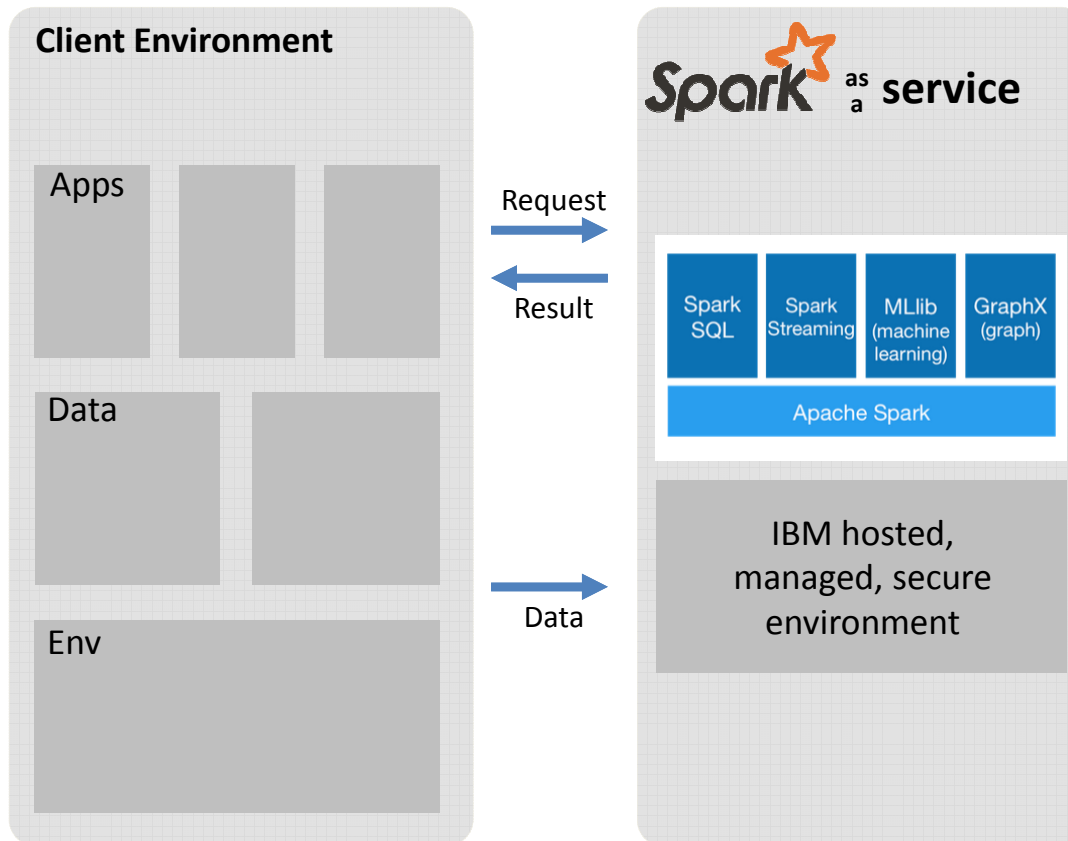
Analytics

Commerce

Watson Health

...many others  
underway

# IBM's Analytics for Apache Spark offering



## What it is:

- Fully-managed Spark environment accessible on-demand for interactive and batch workload

## What you get:

- Access to Spark's next-generation performance and capabilities, including built-in machine learning and others
- Pay only for what you use in either a pay-as-you-go model or through dedicated, enterprise instances
- No lock-in – 100% standard Spark
- Elastic scaling – start with experimentation, extend to development and scale to production, all within the same environment
- Quick start – service is immediately ready for analysis, skipping setup hurdles, hassles and time
- Peace of mind – fully managed and secured, no DBAs or other admins necessary



# Common Spark use cases

1

- Interactive querying of very large data sets (e.g. BI)

2

- Running large data processing batch jobs (e.g. nightly ETL from production systems, primary Hadoop use case)

3

- Complex analytics and data mining across various types of data

4

- Building and deploying rich analytics models (e.g. risk metrics)

5

- Implementing near-realtime stream event processing (e.g. fraud / security detection)

# Spark-as-a-Service targets 4 key personas

## Data Scientist

- Access powerful tools to tease out the insights they're looking for, then make them actionable immediately

```
square <- function(x)
{
  return(x*x)
}
...
```

## Business Analyst

- Answer the questions that the organization needs quickly and easily, and without getting IT involved

```
SELECT
  FirstName, LastName
FROM
  employees
WHERE
  id = 1234;
```

## App Developer

- Add intelligence to their apps in a simple and straightforward no-hassle manner

```
{
  "employees" : {
    "ID" : "1234"
    ...
  }
}
```

## Data Engineer

- Easily build data pipelines that power dashboards and data platforms while ensuring high quality

```
# mysqlimport -u root -
ptmppassword --local test
employee.txt
```



# Challenges in managing Spark applications

- Creating infrastructure silos to accommodate apps is inefficient

Low Utilization  
= Higher cost

*Many new solution workloads in addition to existing apps*



*Leads to costly, complex, siloed, under-utilized infrastructure and replicated data*

## Multi-User/Multi-Tenancy Support

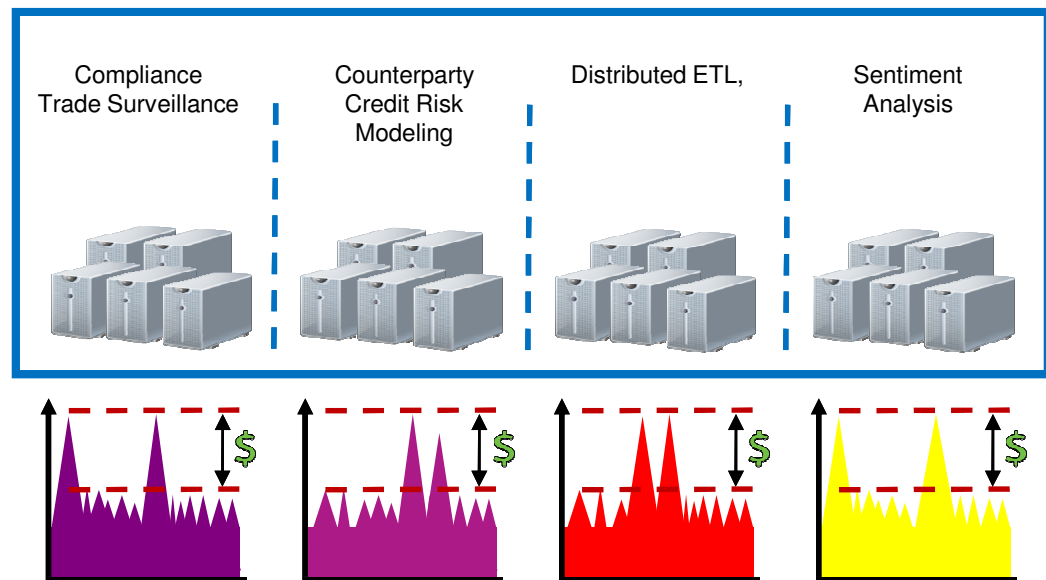
- Different LOBs
- Siloed organization
- Application SLA
- DEV, UAT, PROD

## Spark Lifecycle Management

- Multiple Spark versions
- Different notebooks and versions
- Different data sources, e.g., HDFS, Cassandra

## Enterprise Production Barriers

- Existing applications
- Security, governance
- Monitoring & Logging
- Limited by technology
- It is new!...



# Multi Tenant Spark Cloud Options

- Spark Stand-alone Cluster in VMs
- Use Resource Manager like Mesos/YARN
- Spark-centric Workload Management

# Challenges in Spark Standalone

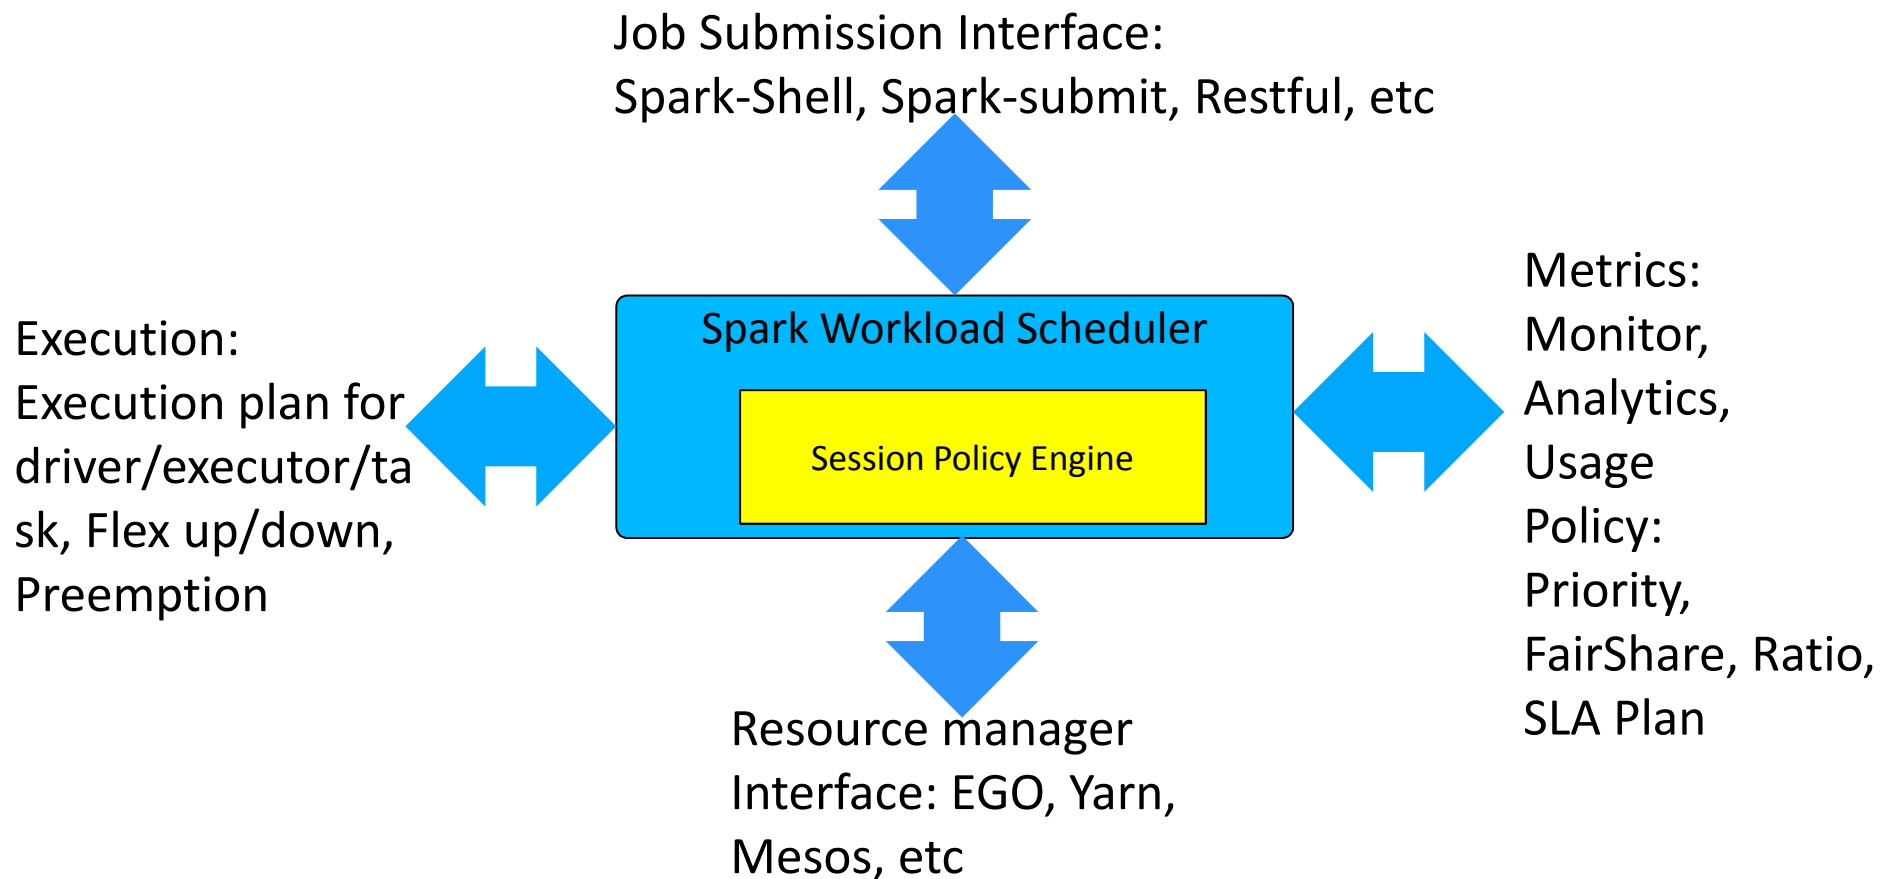


- Adoption – High
  - ✓ 48% user are using standalone deployment
  - ✓ Major cloud provider – Google Dataproc, Databricks, Amazon
- However ...
  - ❖ Single user only – no user grouping for jobs, no authentication, no execution control
  - ❖ No SLA management, application scheduling is not based on real biz priority
  - ❖ Static allocation, rely on user estimate for resource allocation, which needs expertise
  - ❖ No preemption, there could be starvation and deadlock situations
  - ❖ Personal clusters – many silos, high cost in HW and operation/management.
  - ❖ Cannot support multiple frameworks.

# Challenges with Spark on Mesos/YARN

- Adoption – Low
  - ✓ 11% user are using Mesos
  - ✓ Mesos intends to to share the resources across hybrid workloads
- However ...
  - ❖ Fine grained scheduling not really working, no longer supported – Cannot effectively support notebook interactive use case
  - ❖ Single user is running ok, multi-user support requires workload management in Spark scheduler
  - ❖ No SLA management, job round trip time is NOT deterministic when resources are constrained vs demand
  - ❖ Offering mechanism needs to consider resource requirement and demand - lead to in-efficiency, starvation (greedy job keep entire cluster), low utilization.

# Need: Workload Management for Spark Environments



# Agenda

1

Bluemix Spark Cloud Technical  
Challenges

2

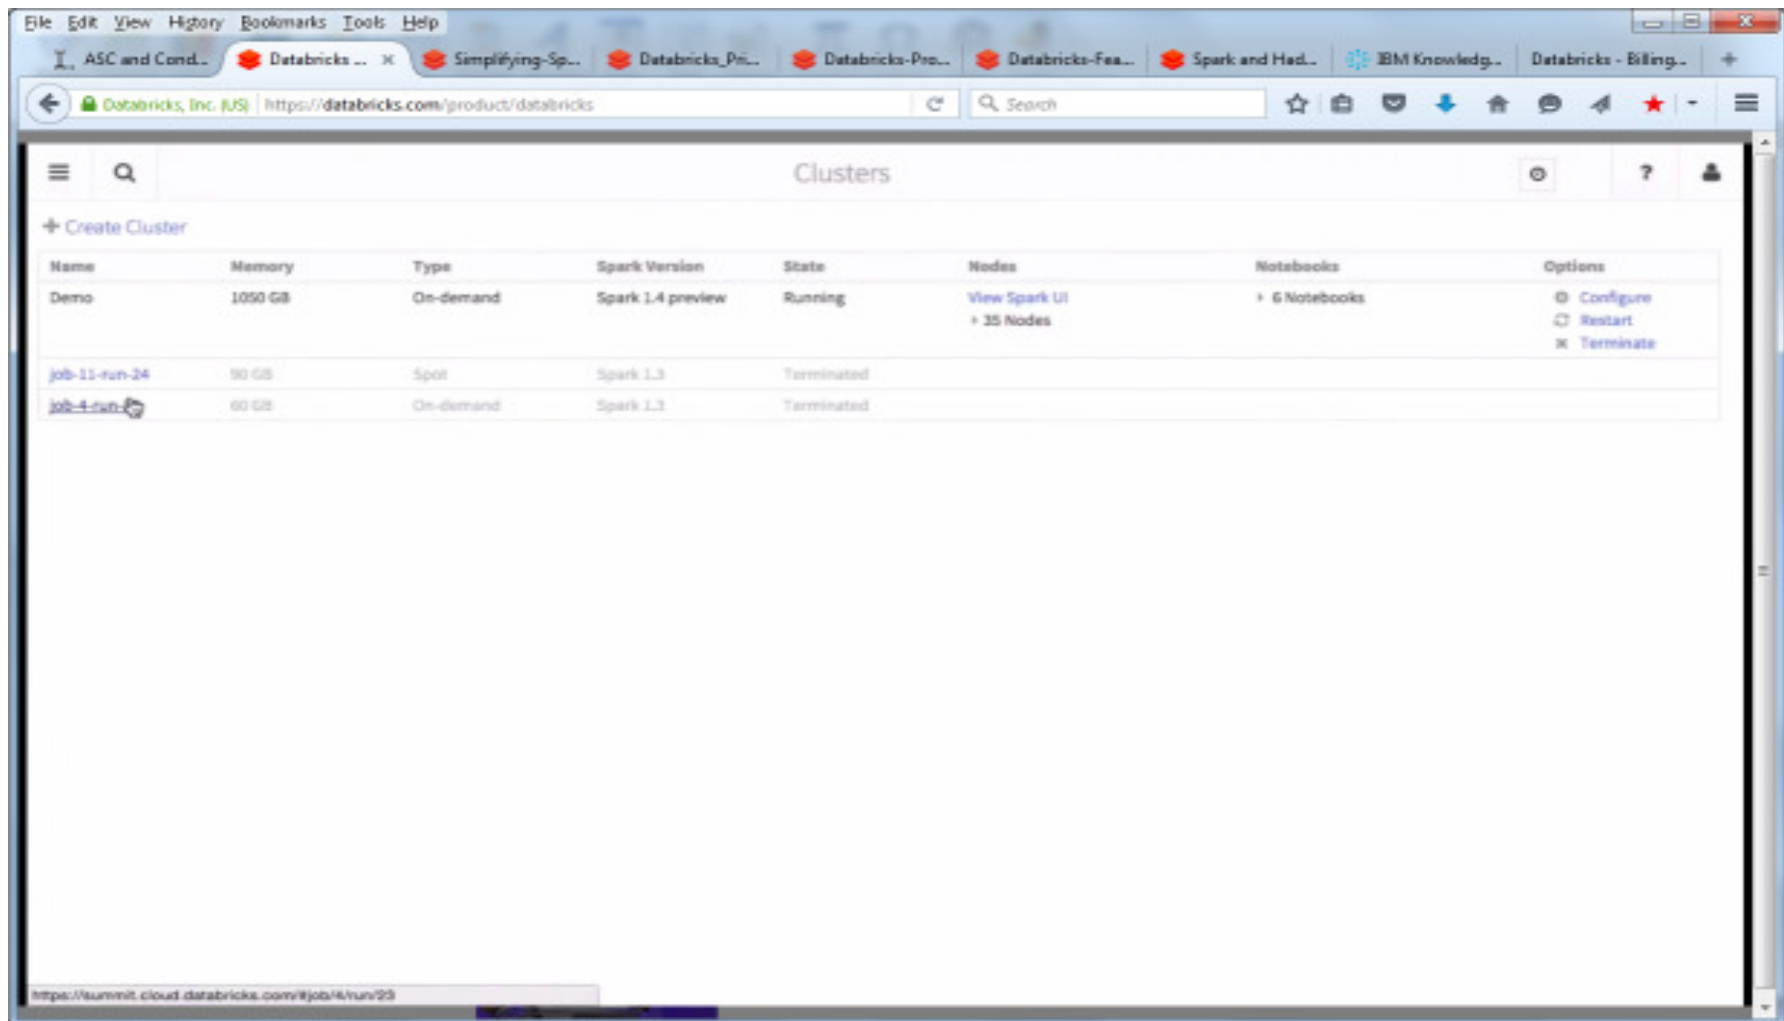
Bluemix Spark Cloud Architecture &  
Implementation

3

IBM – Spark & Mesos Community



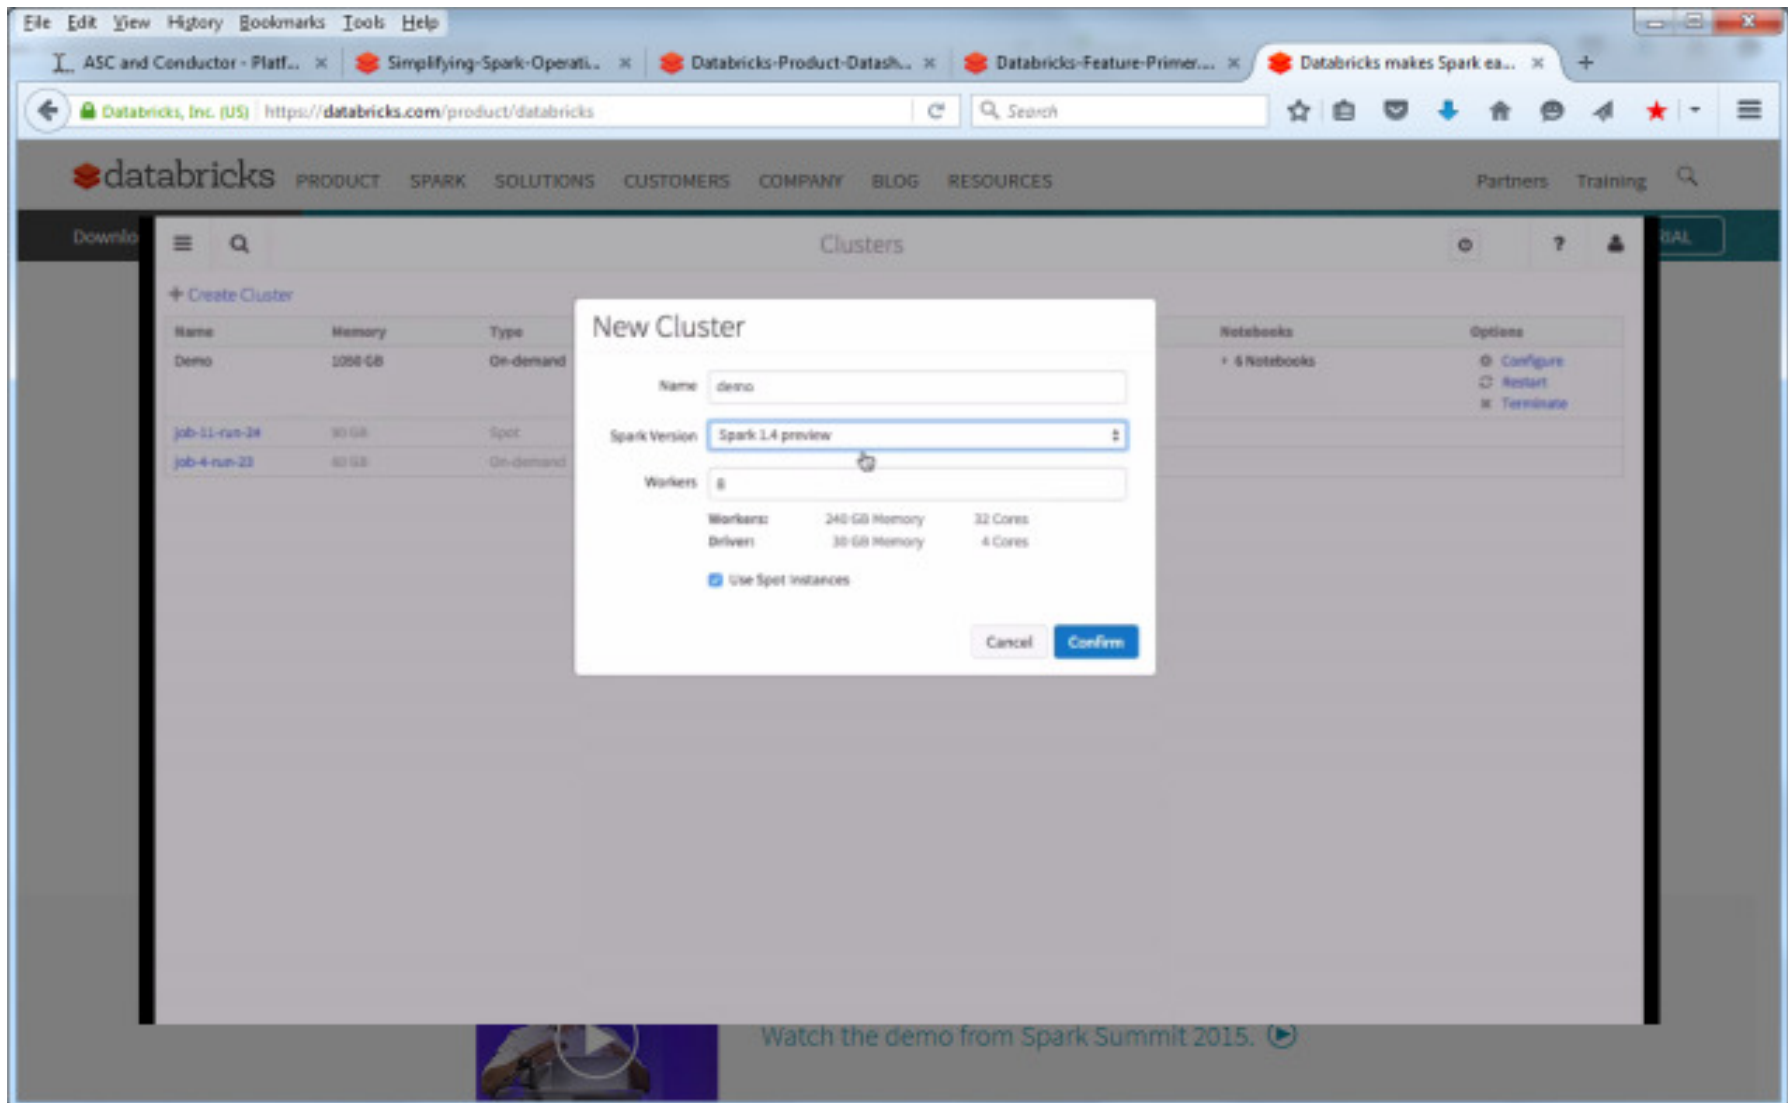
# Spark cloud approach: Spark Clusters



The screenshot shows the Databricks Clusters management interface. The page title is "Clusters". There is a "Create Cluster" button. Below it is a table listing clusters. The table has columns: Name, Memory, Type, Spark Version, State, Nodes, Notebooks, and Options. The first cluster, "Demo", is in a "Running" state with 1050 GB of memory, On-demand type, Spark 1.4 preview version, and 35 nodes. The other two clusters, "job-11-run-24" and "job-4-run-23", are in a "Terminated" state. The "job-4-run-23" cluster is highlighted with a mouse cursor. The "Options" column for the "Demo" cluster includes links for "Configure", "Restart", and "Terminate".

Name	Memory	Type	Spark Version	State	Nodes	Notebooks	Options
Demo	1050 GB	On-demand	Spark 1.4 preview	Running	<a href="#">View Spark UI</a> + 35 Nodes	+ 6 Notebooks	<a href="#">Configure</a> <a href="#">Restart</a> <a href="#">Terminate</a>
job-11-run-24	90 GB	Spot	Spark 1.3	Terminated			
job-4-run-23	60 GB	On-demand	Spark 1.3	Terminated			

## Spark Cloud approach: New Spark Cluster per tenant



[Back to All Categories](#)

## Apache Spark

IBM

PUBLISH DATE  
07/05/2016AUTHOR  
IBMTYPE  
Service[VIEW DOCS](#)

Apache Spark is an open source cluster computing framework optimized for extremely fast and large scale data processing, which you can access via the newly integrated notebook interface IBM Analytics for Apache Spark. You can connect to your existing data sources or take advantage of the on-demand big data optimization of Object Storage. Spark plans are based on the maximum number of executors available to process your analytic jobs. Executors exist only as long as they're needed for processing, so you're charged only for processing done.

- **Incredibly Fast**

Apache Spark delivers 100x the performance of Apache Hadoop for certain workloads because of its advanced in-memory computing engine.

- **Easy to Use and Powerful**

Apache Spark's Streaming and SQL programming models backed by MLlib and GraphX make it incredibly easy for developers and data scientists to build apps that exploit machine learning and graph analytics. Because the service is 100% compatible with Apache Spark, developers can build their apps and run them against the IBM managed service to benefit from operational, maintenance, and hardware excellence.

- **Convenient Data Storage**

Object Storage enables a convenient way to upload your data from a file for immediate use by your Spark instance. You can set up Object Storage directly from the Spark service interface.

### Pick a plan

Monthly prices shown are for country or region: [United States](#)

Plan	Features	
<a href="#">Personal</a>	2 Spark Executors	\$0.70 USD/Instance-Hour
✓ <a href="#">Reserved Enterprise</a>	30 Spark Executors	-



A plan to run programs using up to 30 Spark executors. Email [sparksrv@us.ibm.com](mailto:sparksrv@us.ibm.com) to order.

Begin composing your service with

**Apache Spark**[LOG IN TO BLUEMIX](#)

Don't have an account?

[SIGN UP FOR A FREE TRIAL](#)

# Jupyter interactive notebooks



Notebooks are **interactive** computational environments, in which you can **combine** code execution, rich text, mathematics, plots and rich media.

## Simple spectral analysis

An illustration of the [Discrete Fourier Transform](#)

$$X_k = \sum_{n=0}^{N-1} x_n e^{-\frac{2\pi i}{N} kn} \quad k = 0, \dots, N-1$$


using windowing, to reveal the frequency content of a sound signal.

We begin by loading a datafile using SciPy's audio file support:

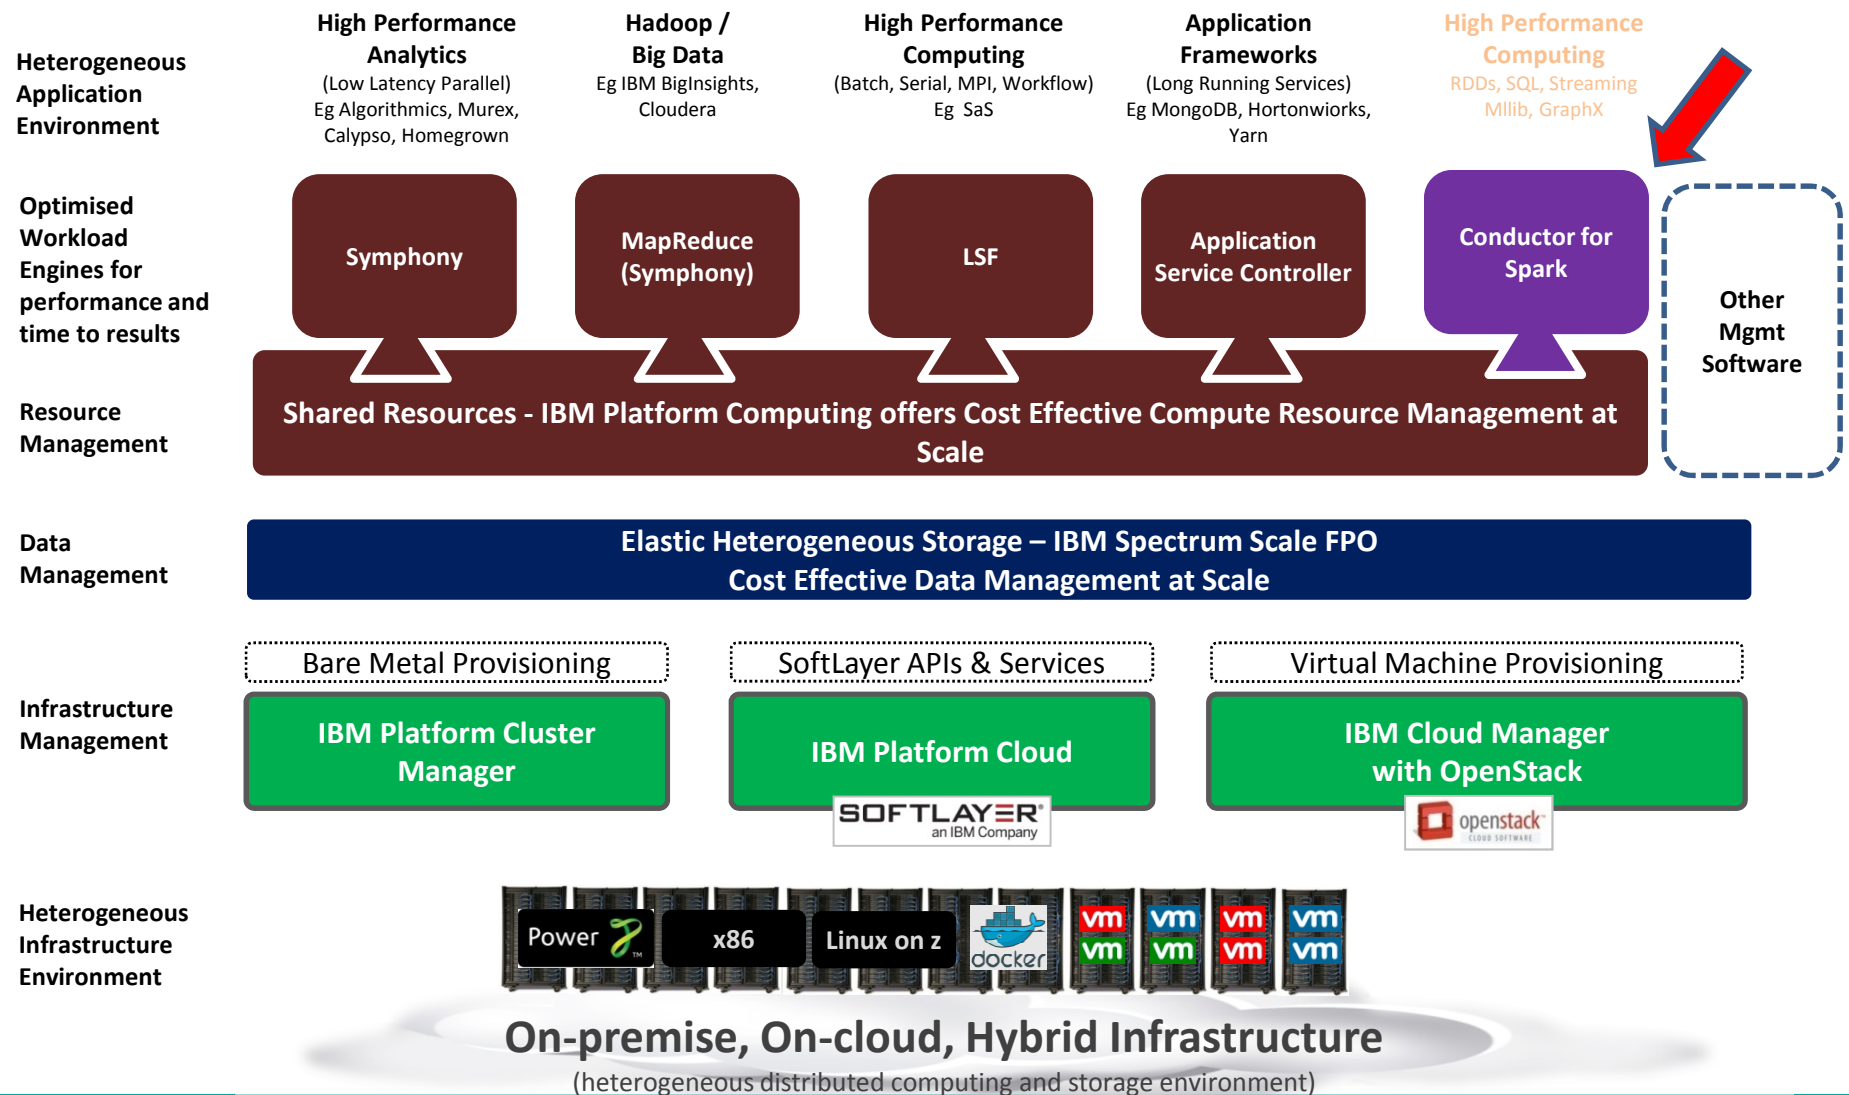
```
In [1]: from scipy.io import wavfile
rate, x = wavfile.read('test_mono.wav')
```

And we can easily view its spectral structure using matplotlib's builtin specgram routine:

```
In [2]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 4))
ax1.plot(x); ax1.set_title('Raw audio signal')
ax2.specgram(x); ax2.set_title('Spectrogram');
```

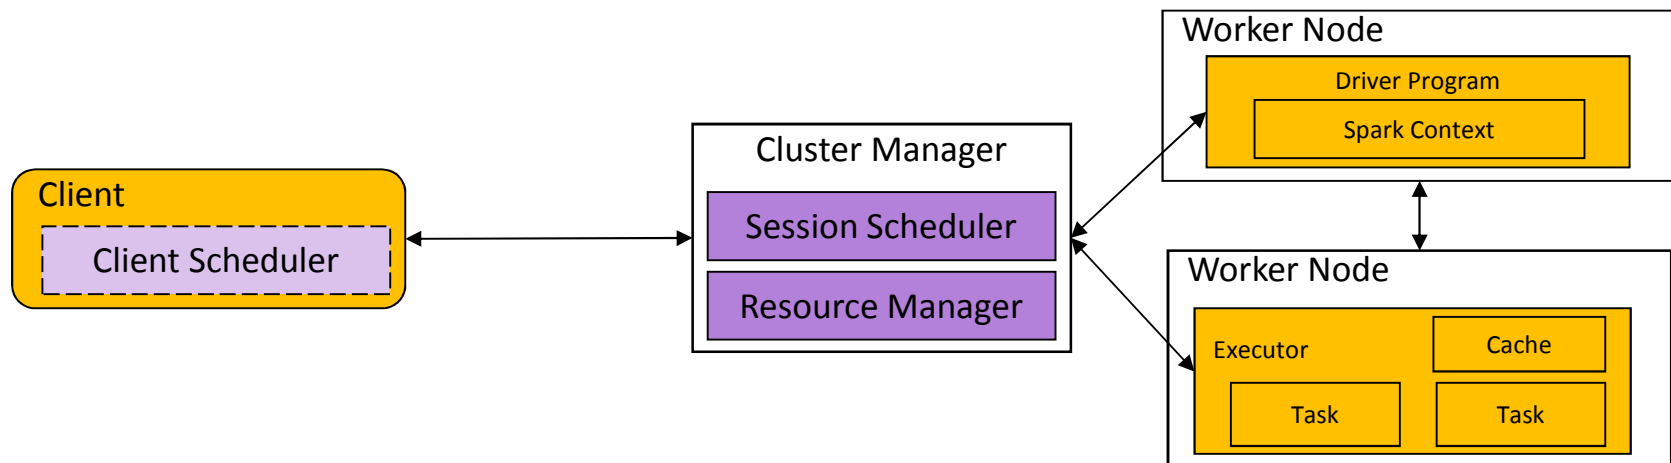


# IBM Software Defined Infrastructure Extends to Conductor for Spark

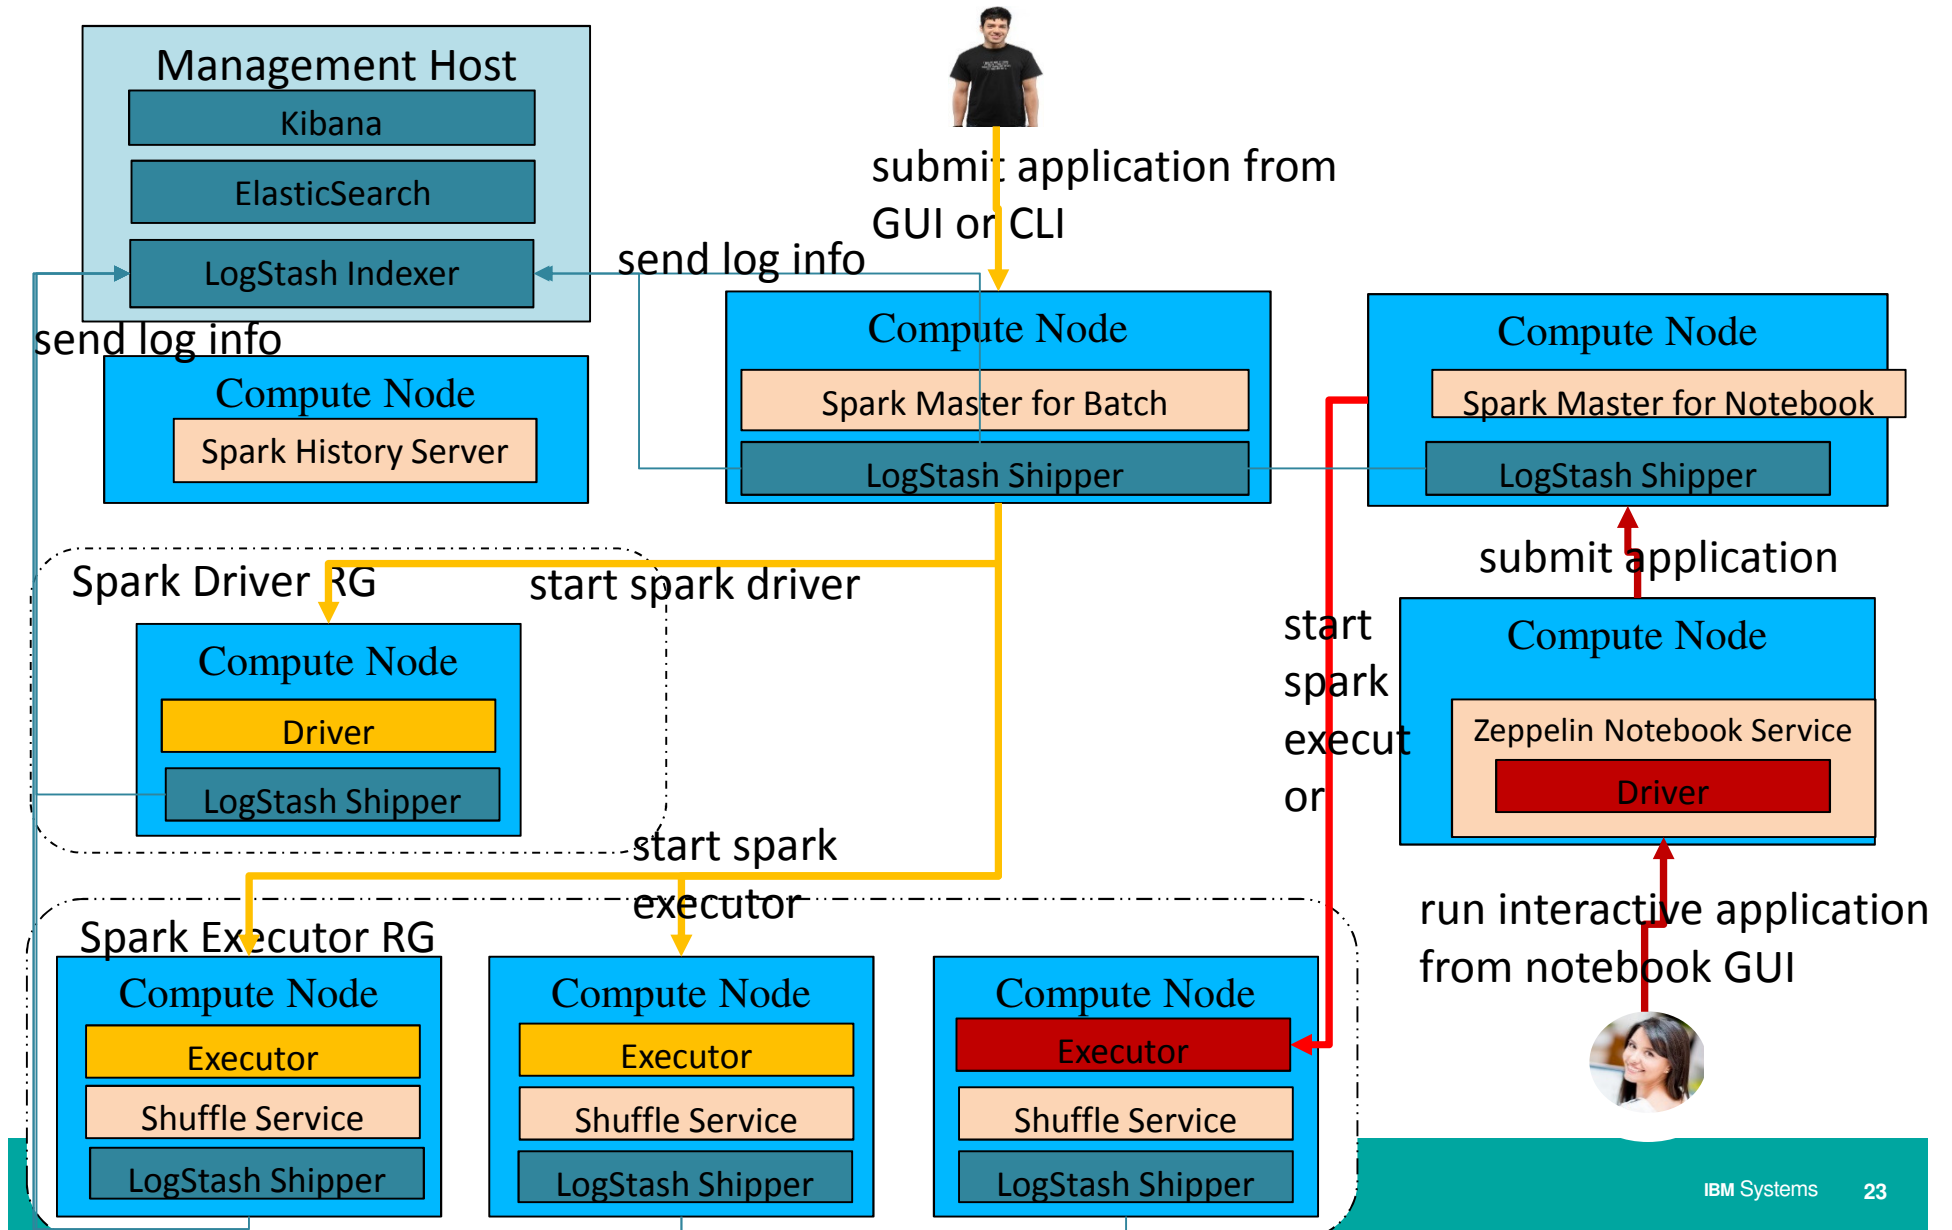


## Workload Scheduling with Session Scheduler

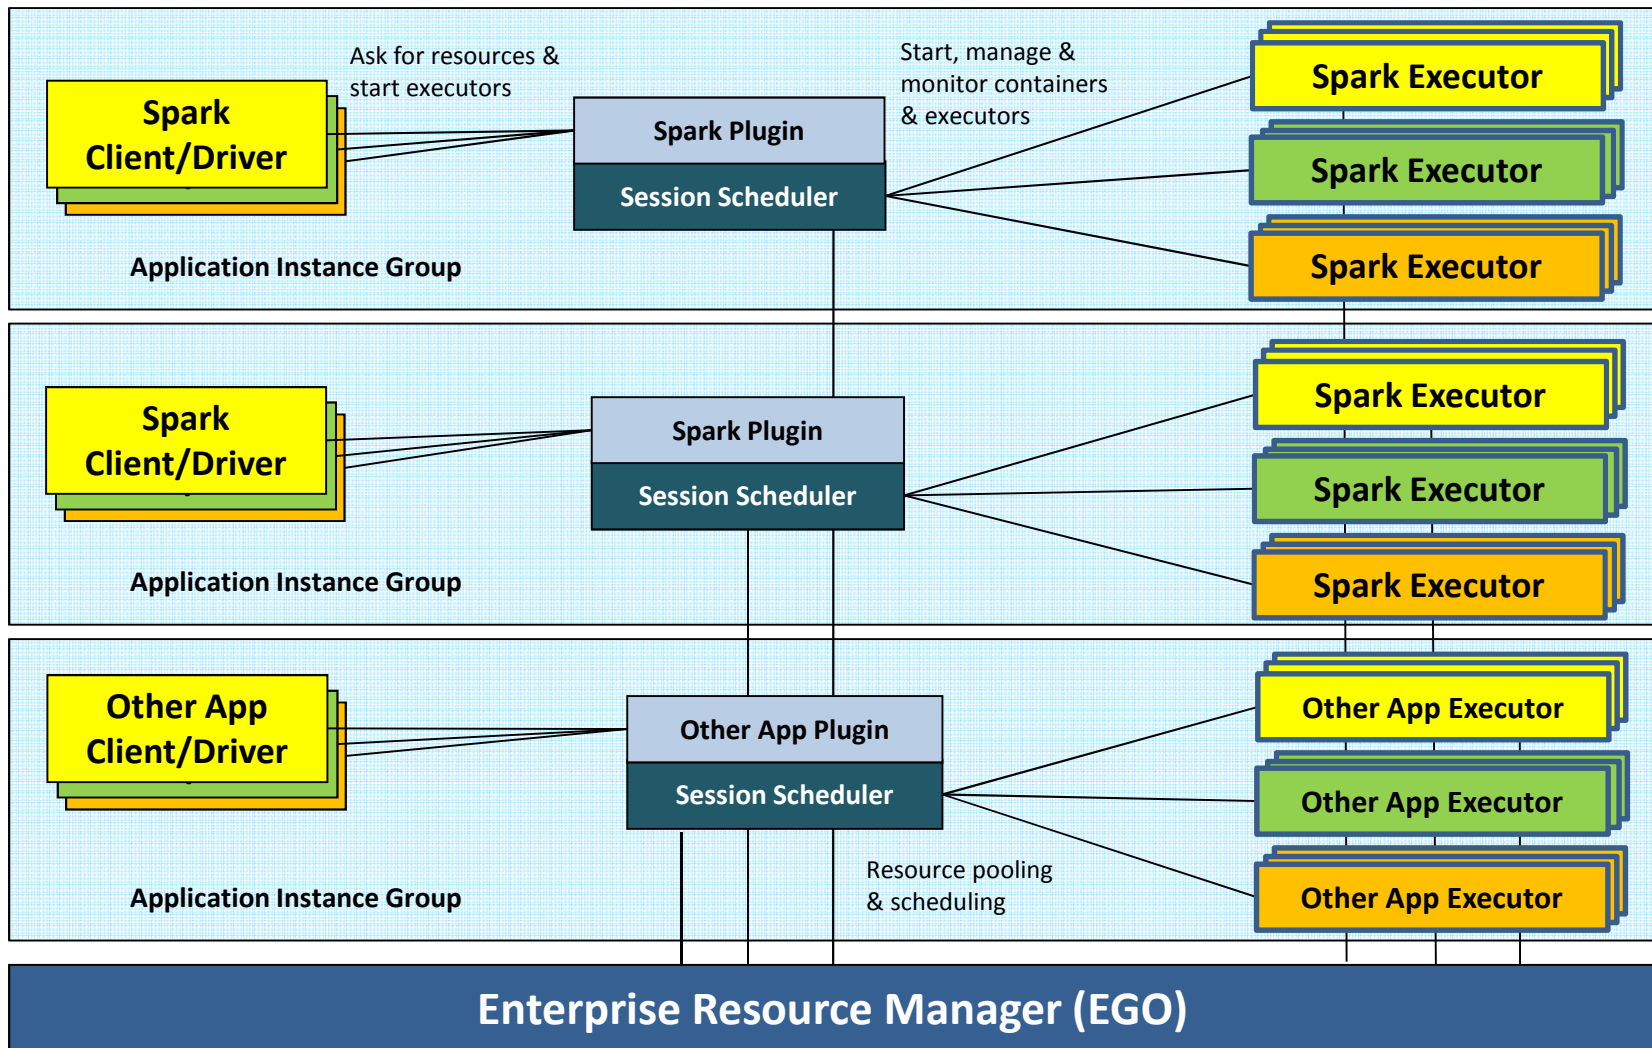
- Improved task scheduling, No changes to Apache Spark APIs
  - Multi-tenant scheduling with Session Scheduler
  - Resource sharing and reclaim
  - Fine-grained scheduling
  - Multiple resource scheduling policy ( FIFO / Fairshare) for each tenant



# Architecture: Runtime of a Spark Instance Group



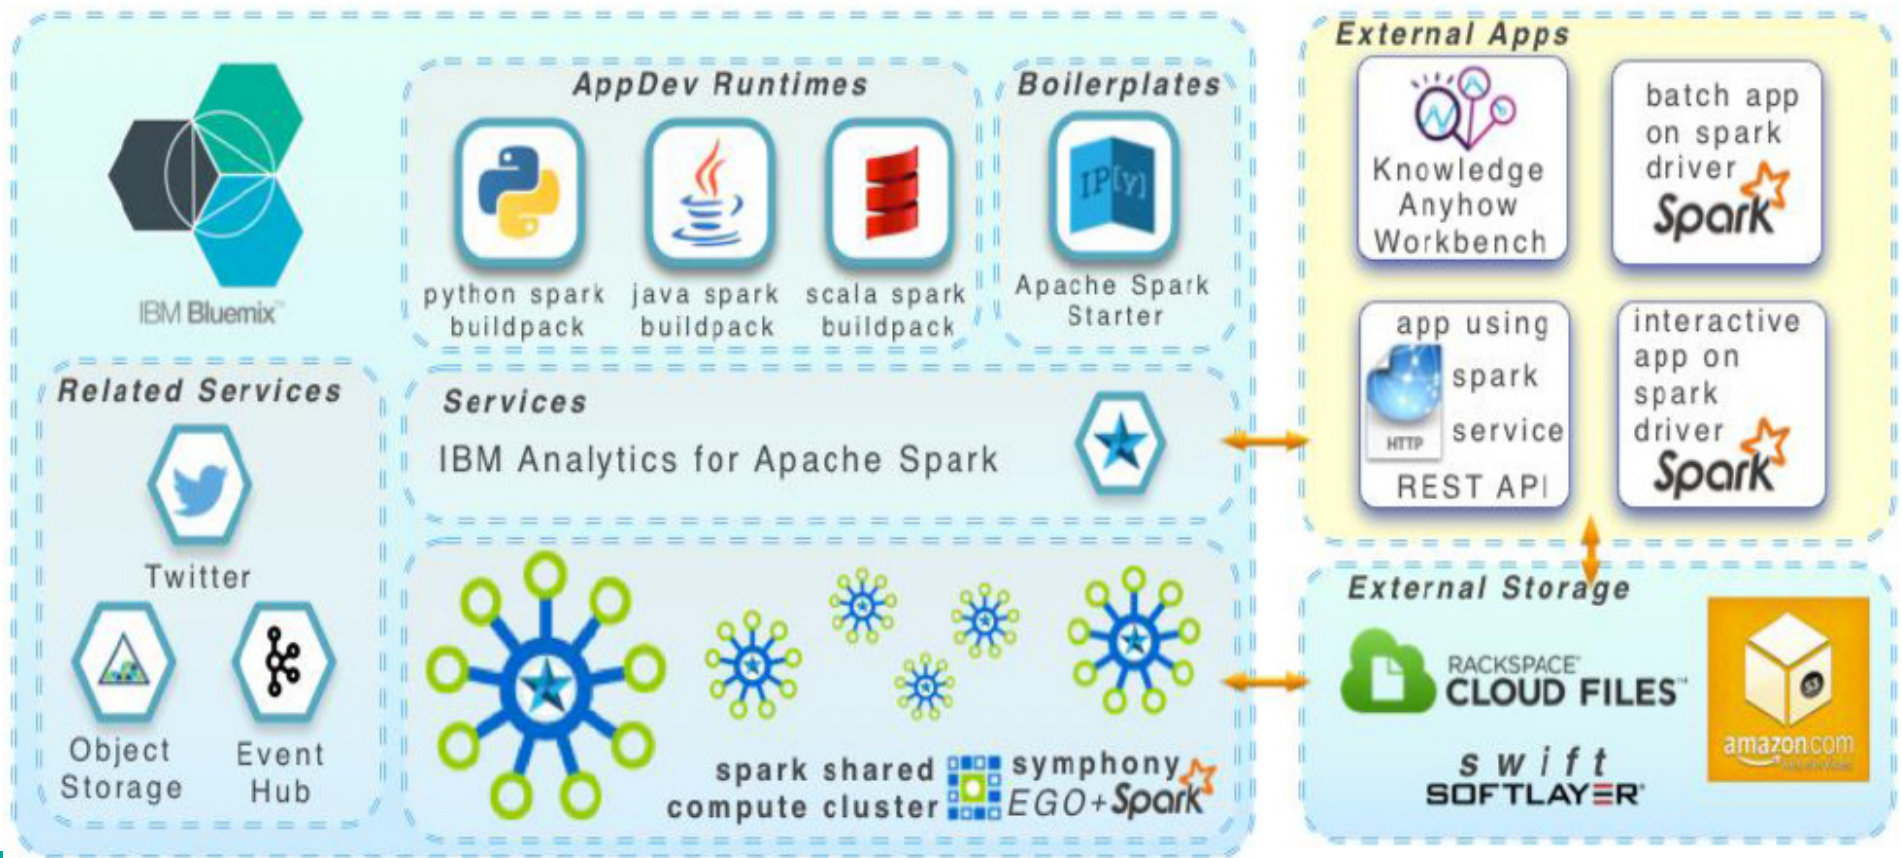
## *Deploy, manage and schedule multiple application instance groups*



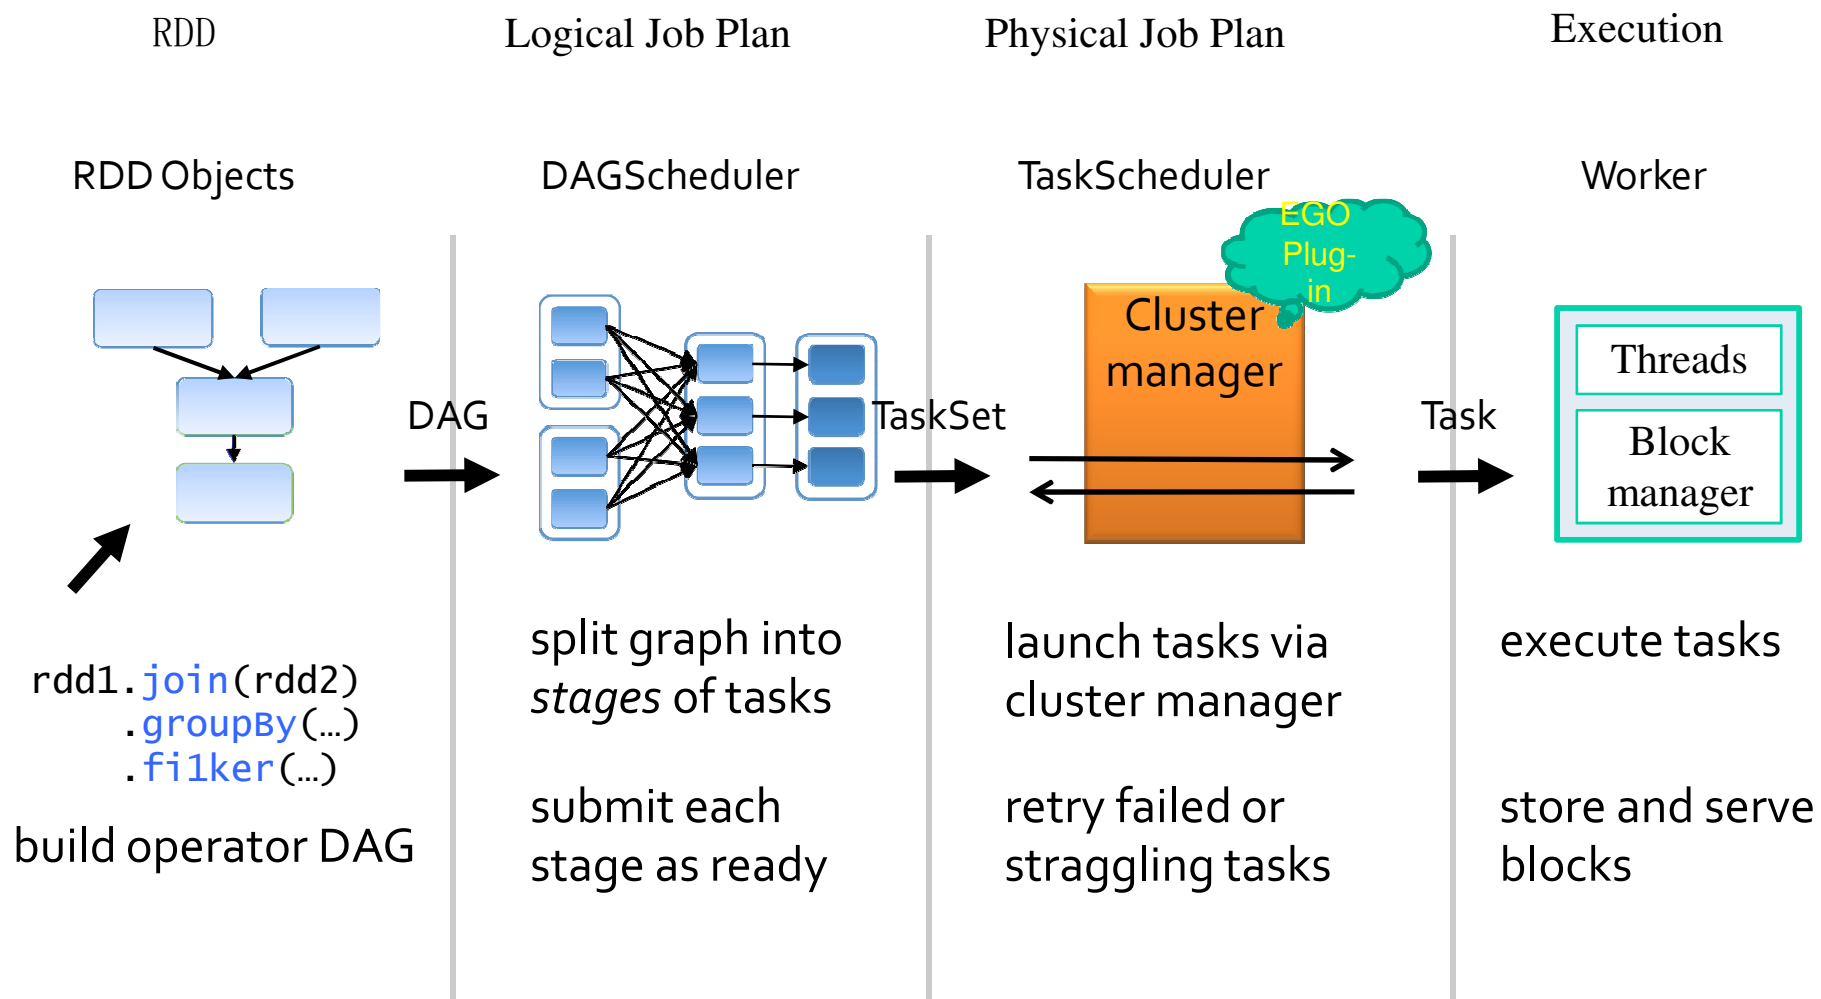


# IBM Analytics Spark Cloud Service

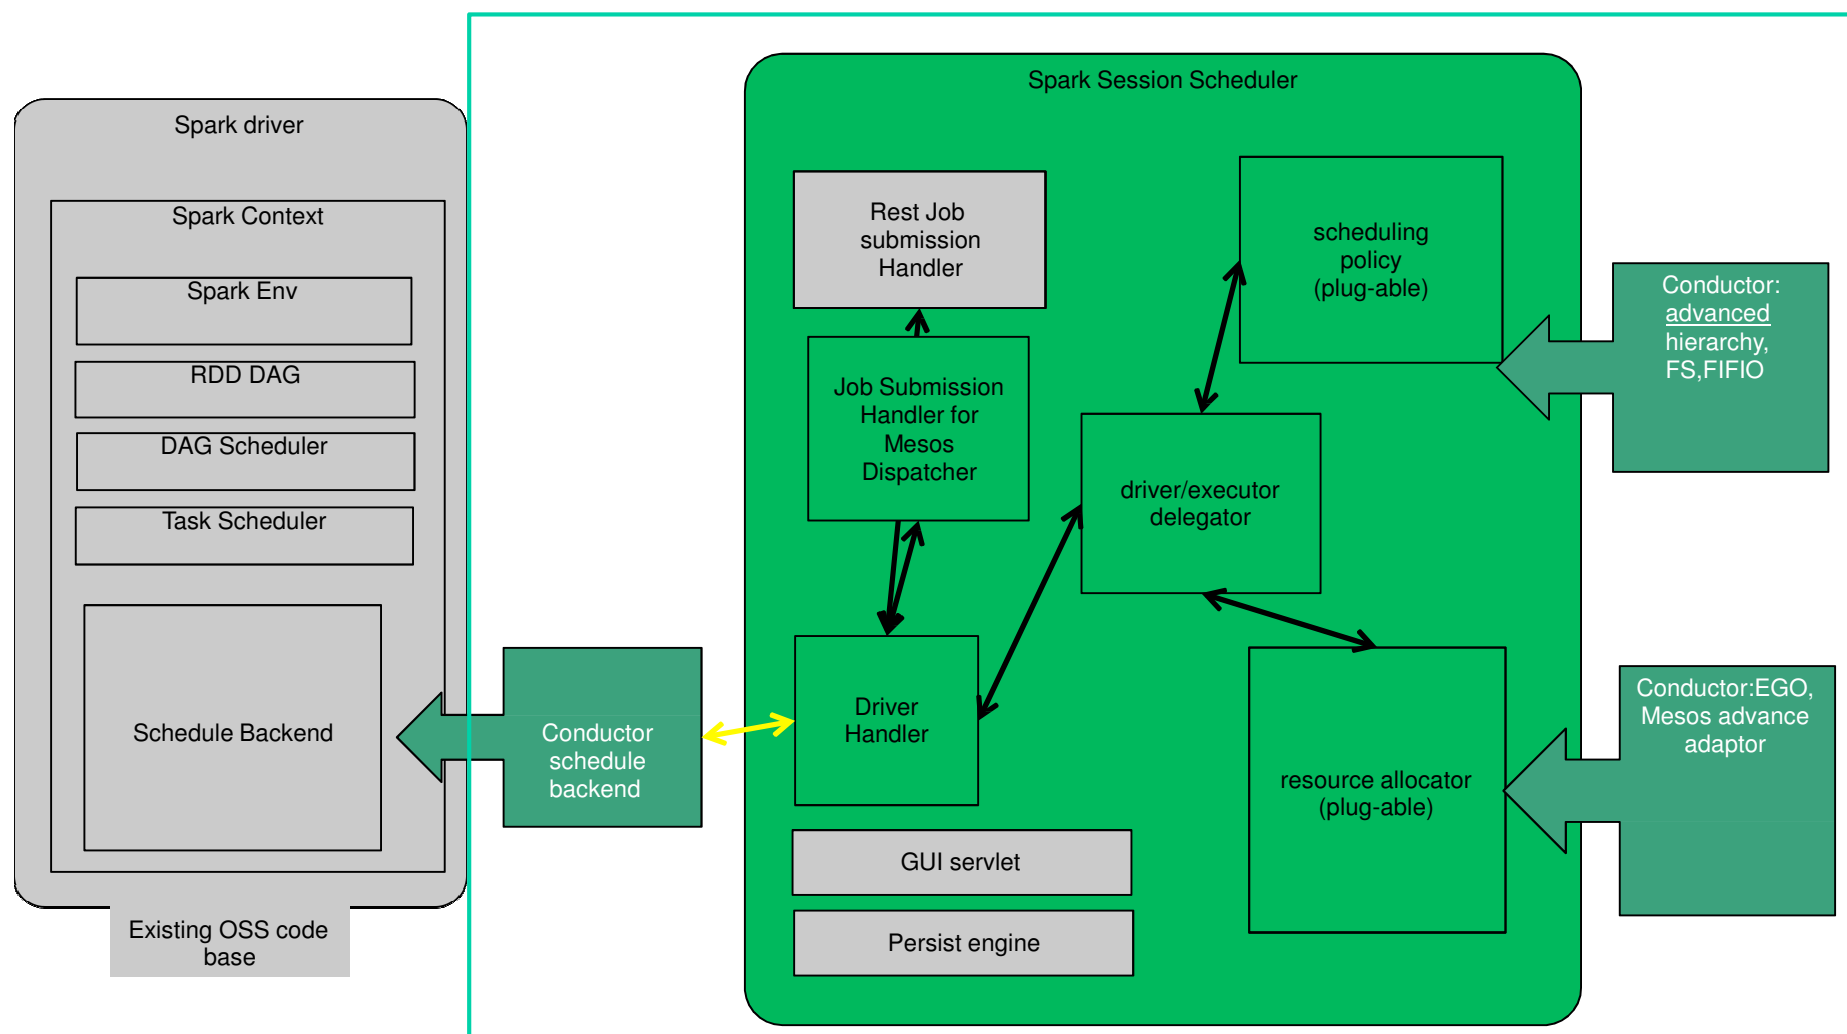
- Multi-Tenancy
- Service Provisioning
- Fine Grain Scheduling
- Security Isolation
- High Resource Utilization
- Enterprise Class Solution



# How Scheduler plug-in works with Apache Spark

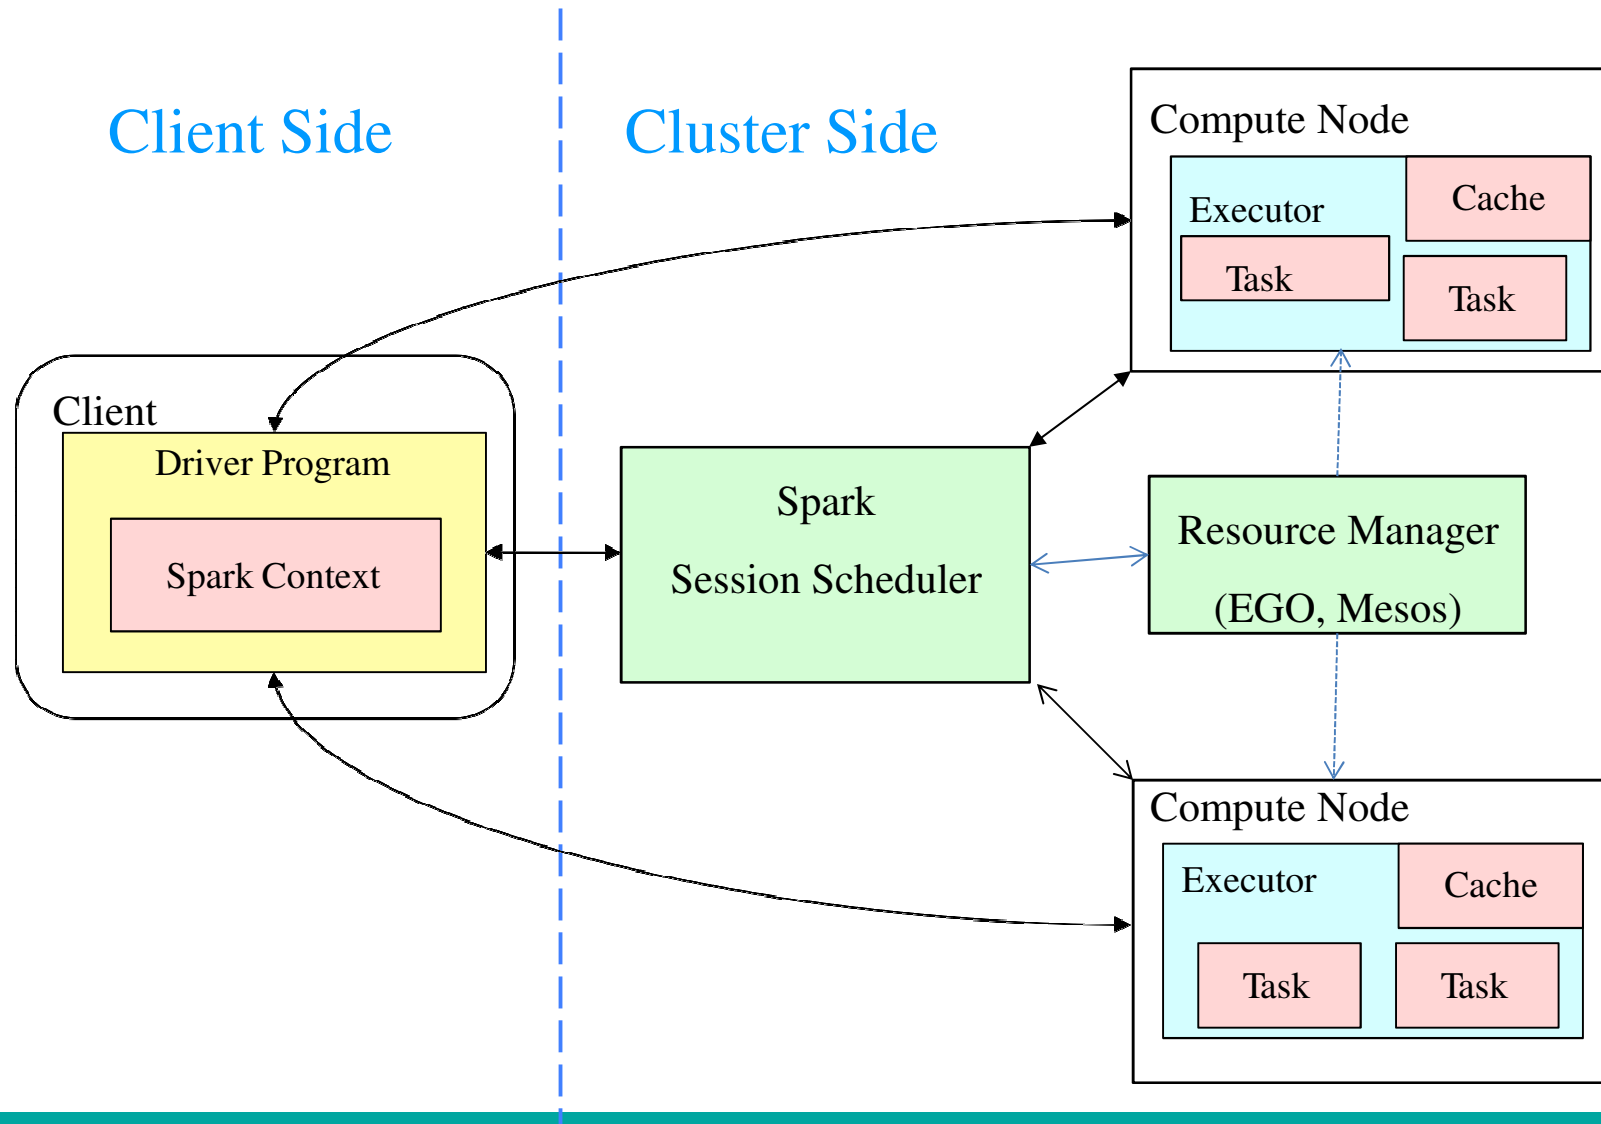


# Advanced Spark Session Scheduler

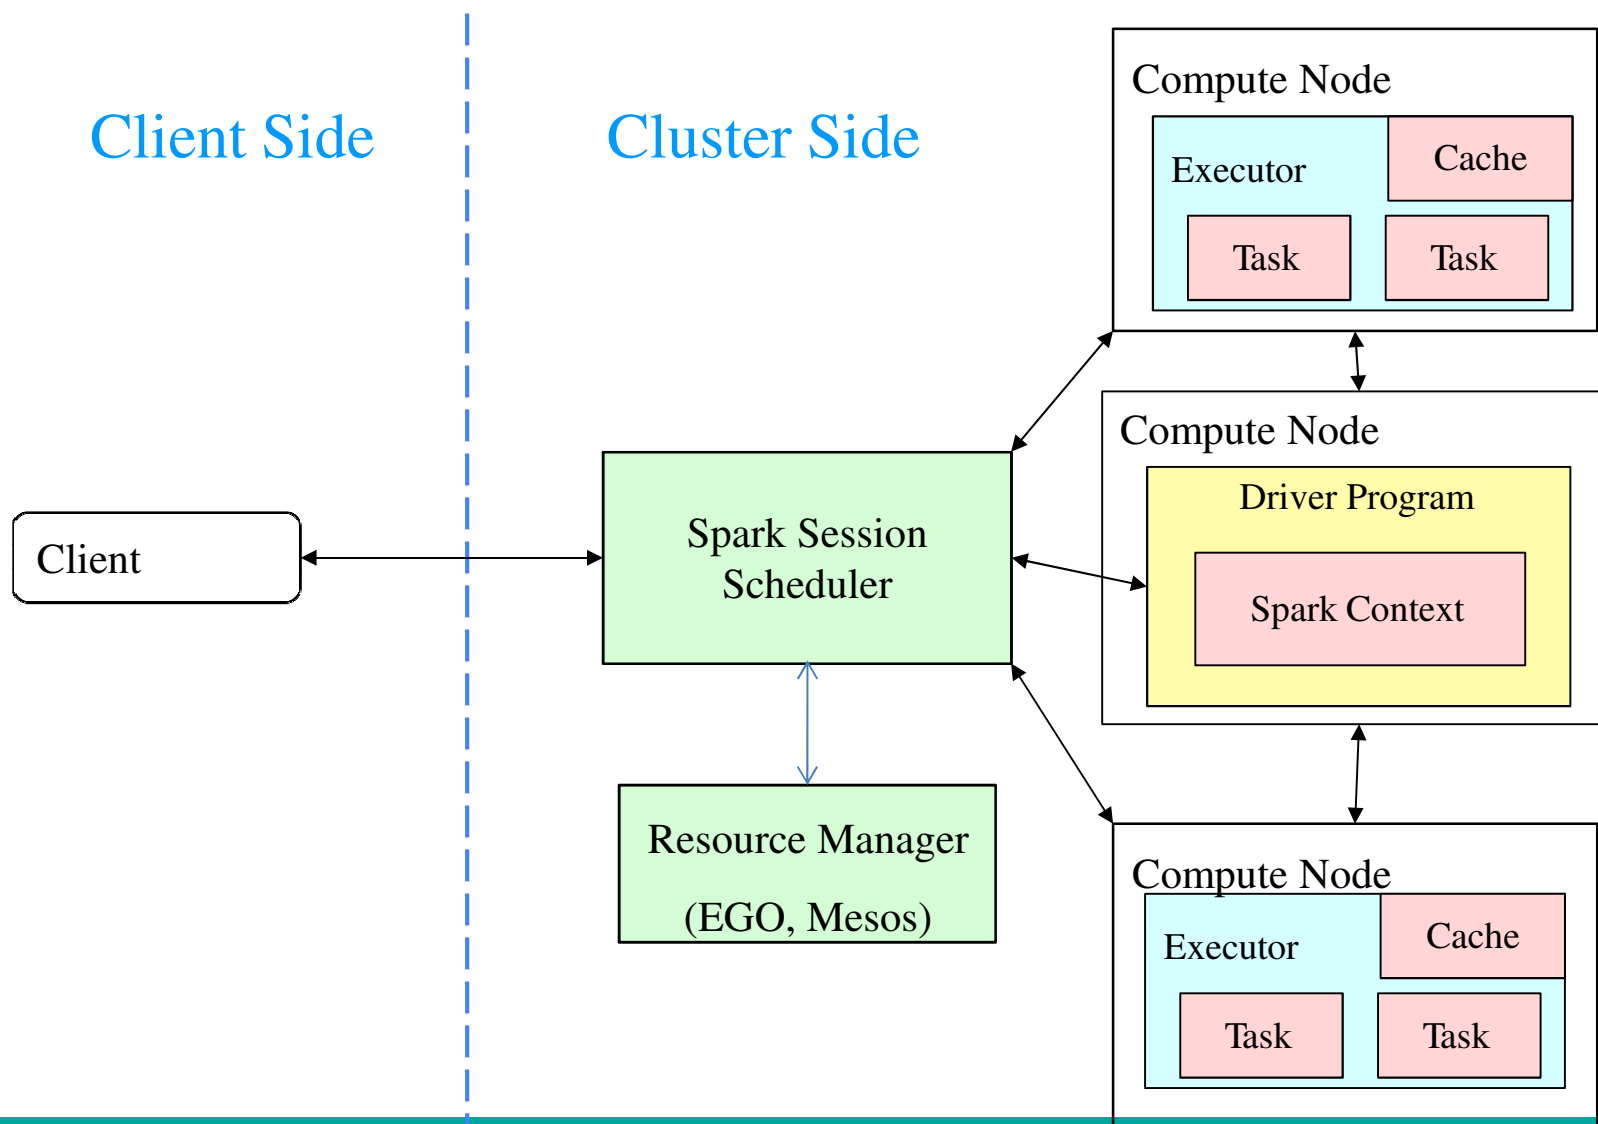


Conductor

# Spark Scheduler - Deploy: client mode / spark-shell





# Spark Scheduler - Deploy: cluster mode



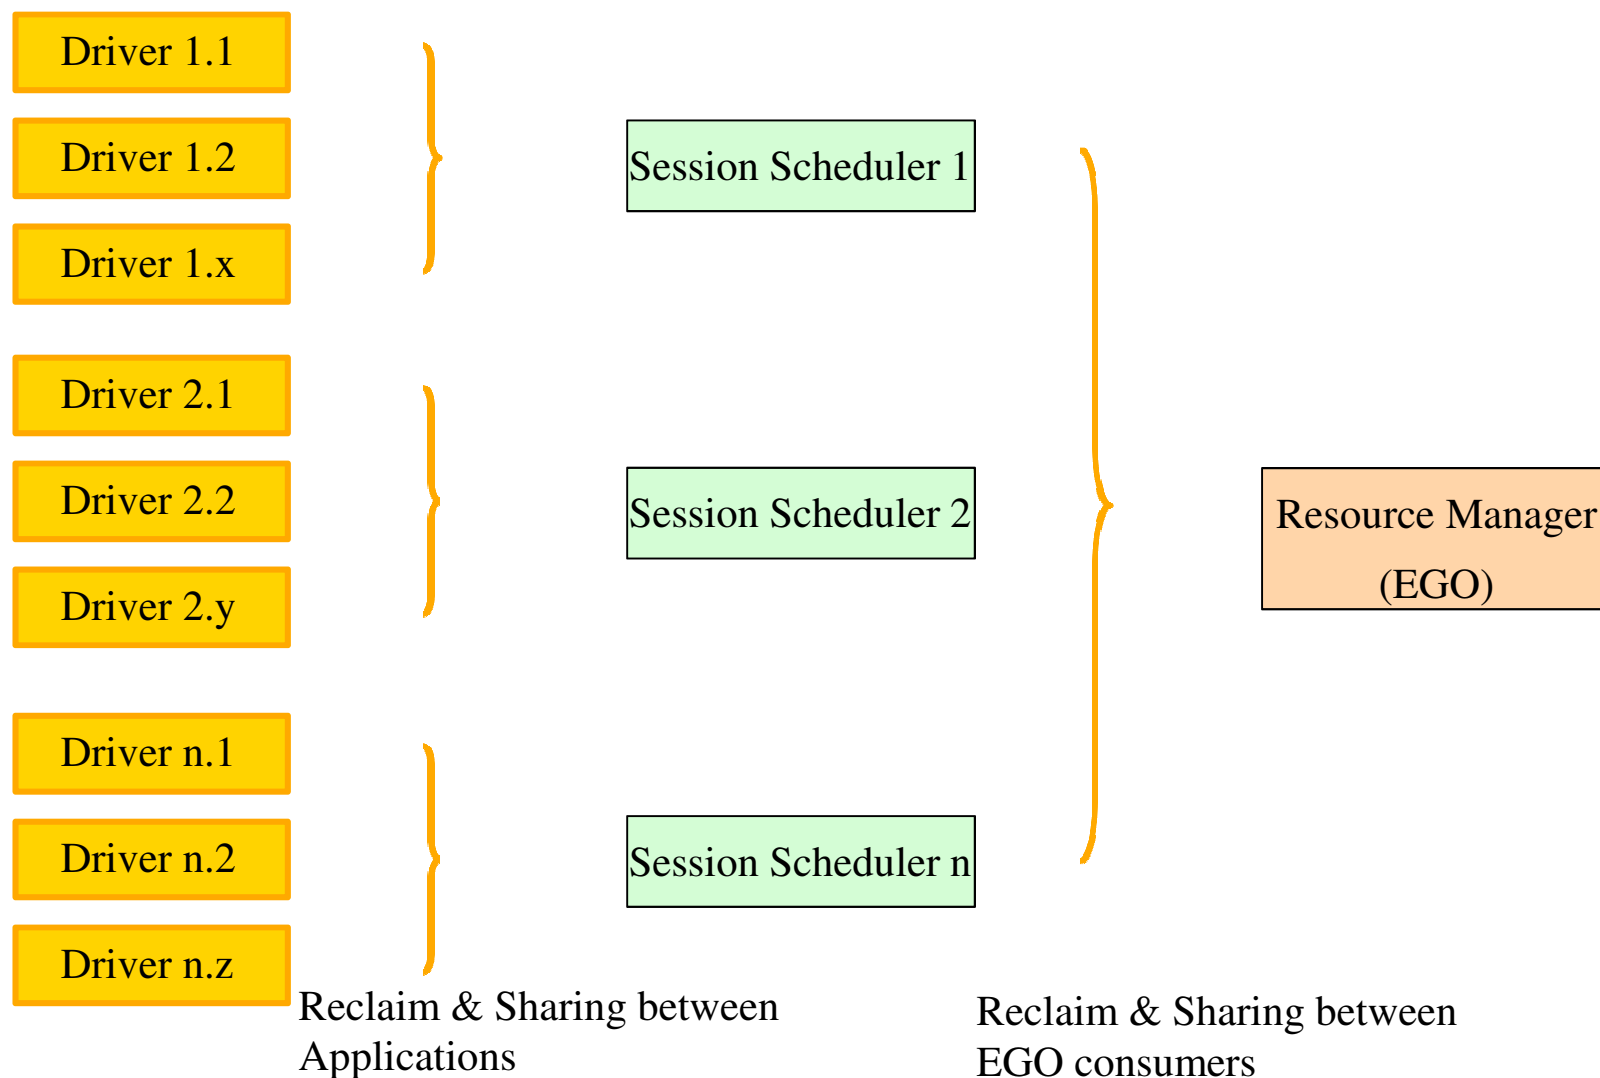
# Spark Session Scheduler Advantages - Fine-grained scheduling

- Opensource Spark (Standalone, YARN, Mesos) assign resources by executor.
- In Conductor with Spark, resources are allocated on demand at task level.
- Fine-grained scheduling also achieves better data locality.

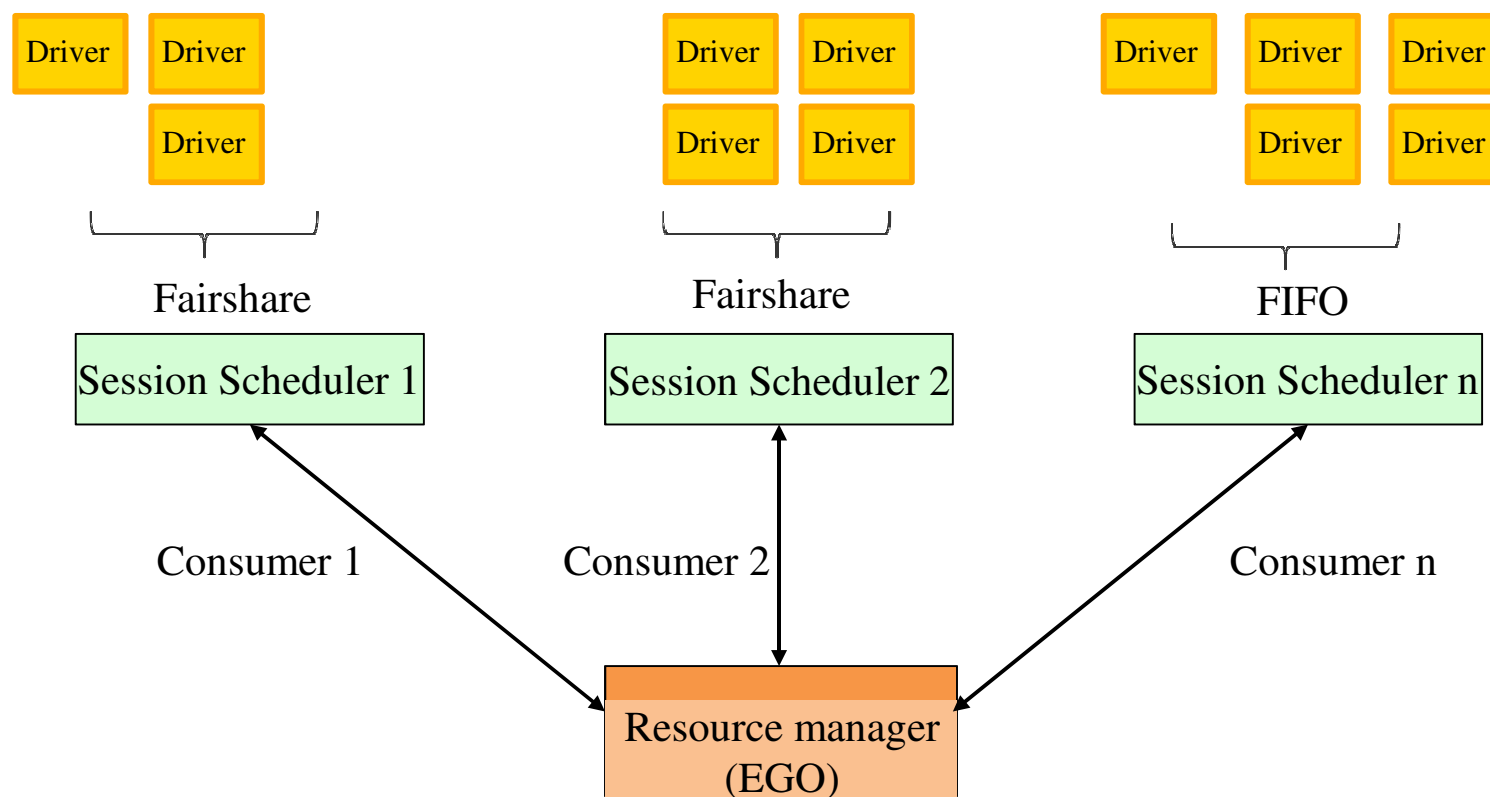
Stage	S1	S2	S3	S4	S5	S6
Task No	20	7	15	3	24	2



# Spark on EGO Advantages - Reclaim & Sharing

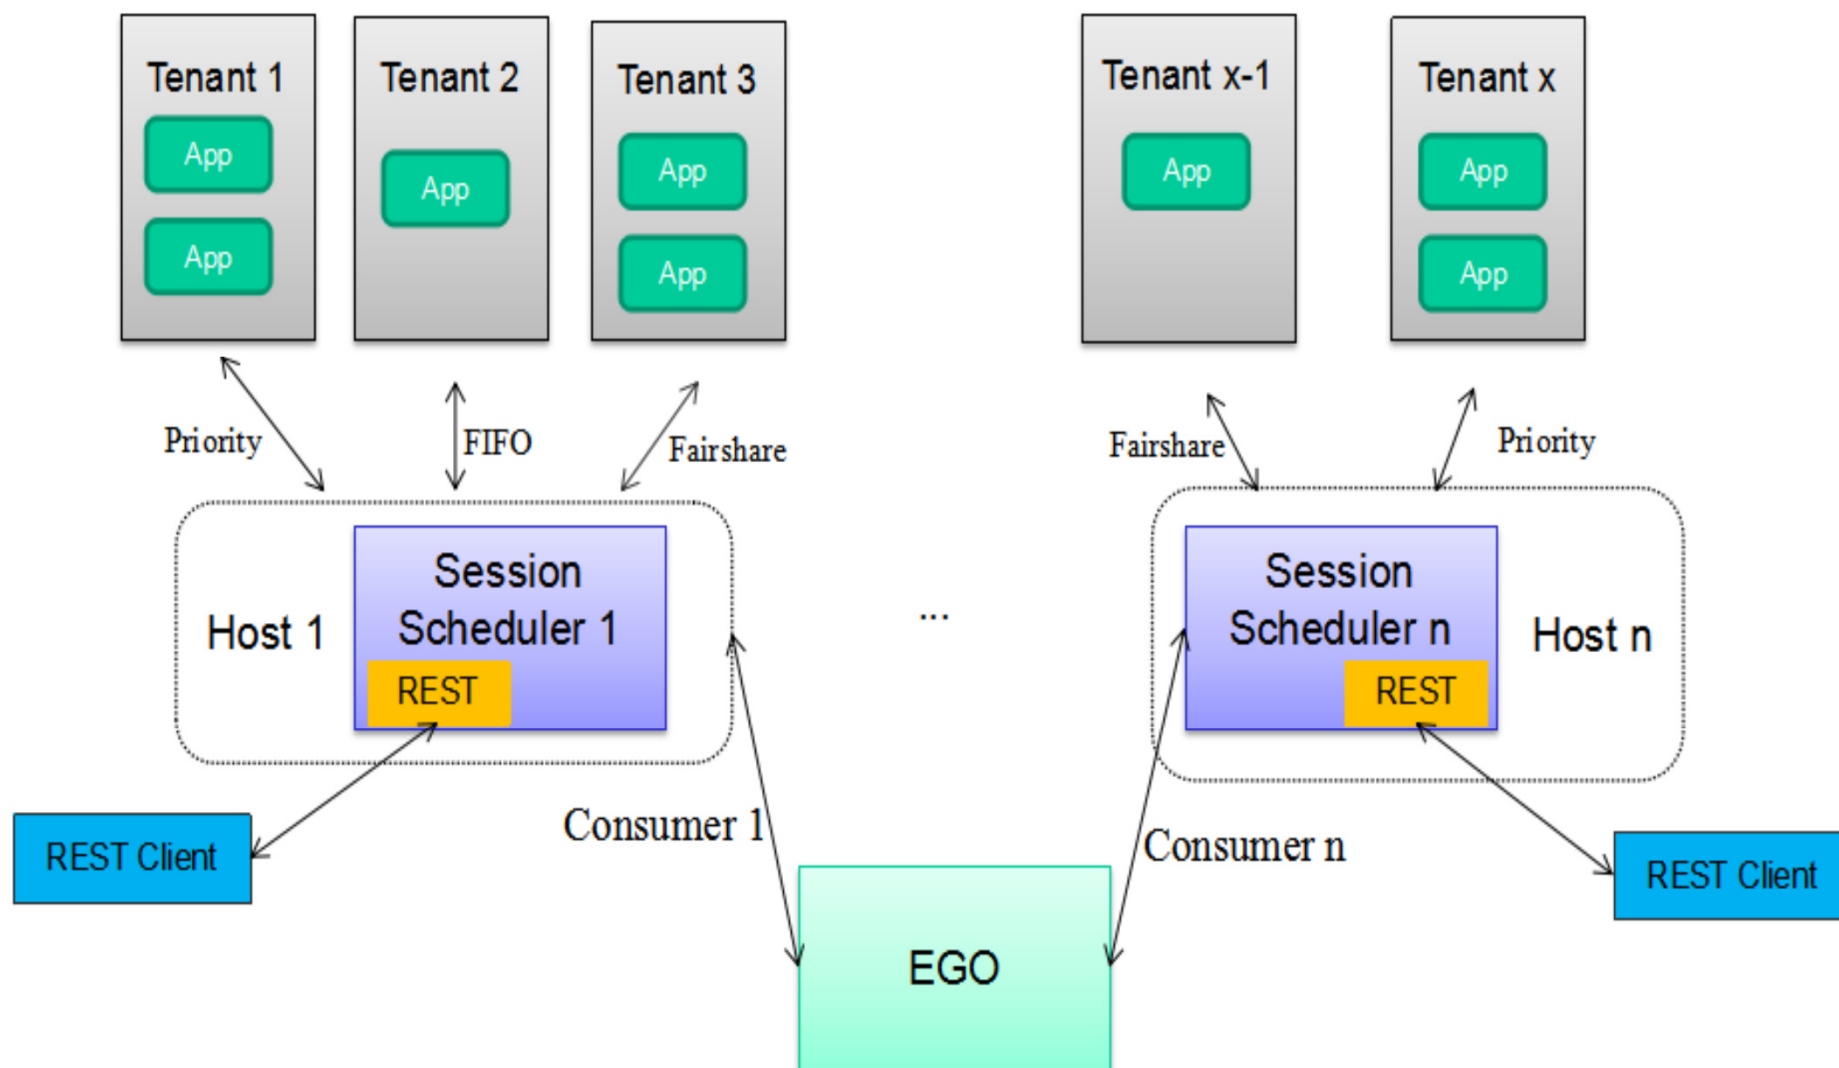


## Spectrum Conductor with Spark: multi-tenant & multiple-resource schedule policy

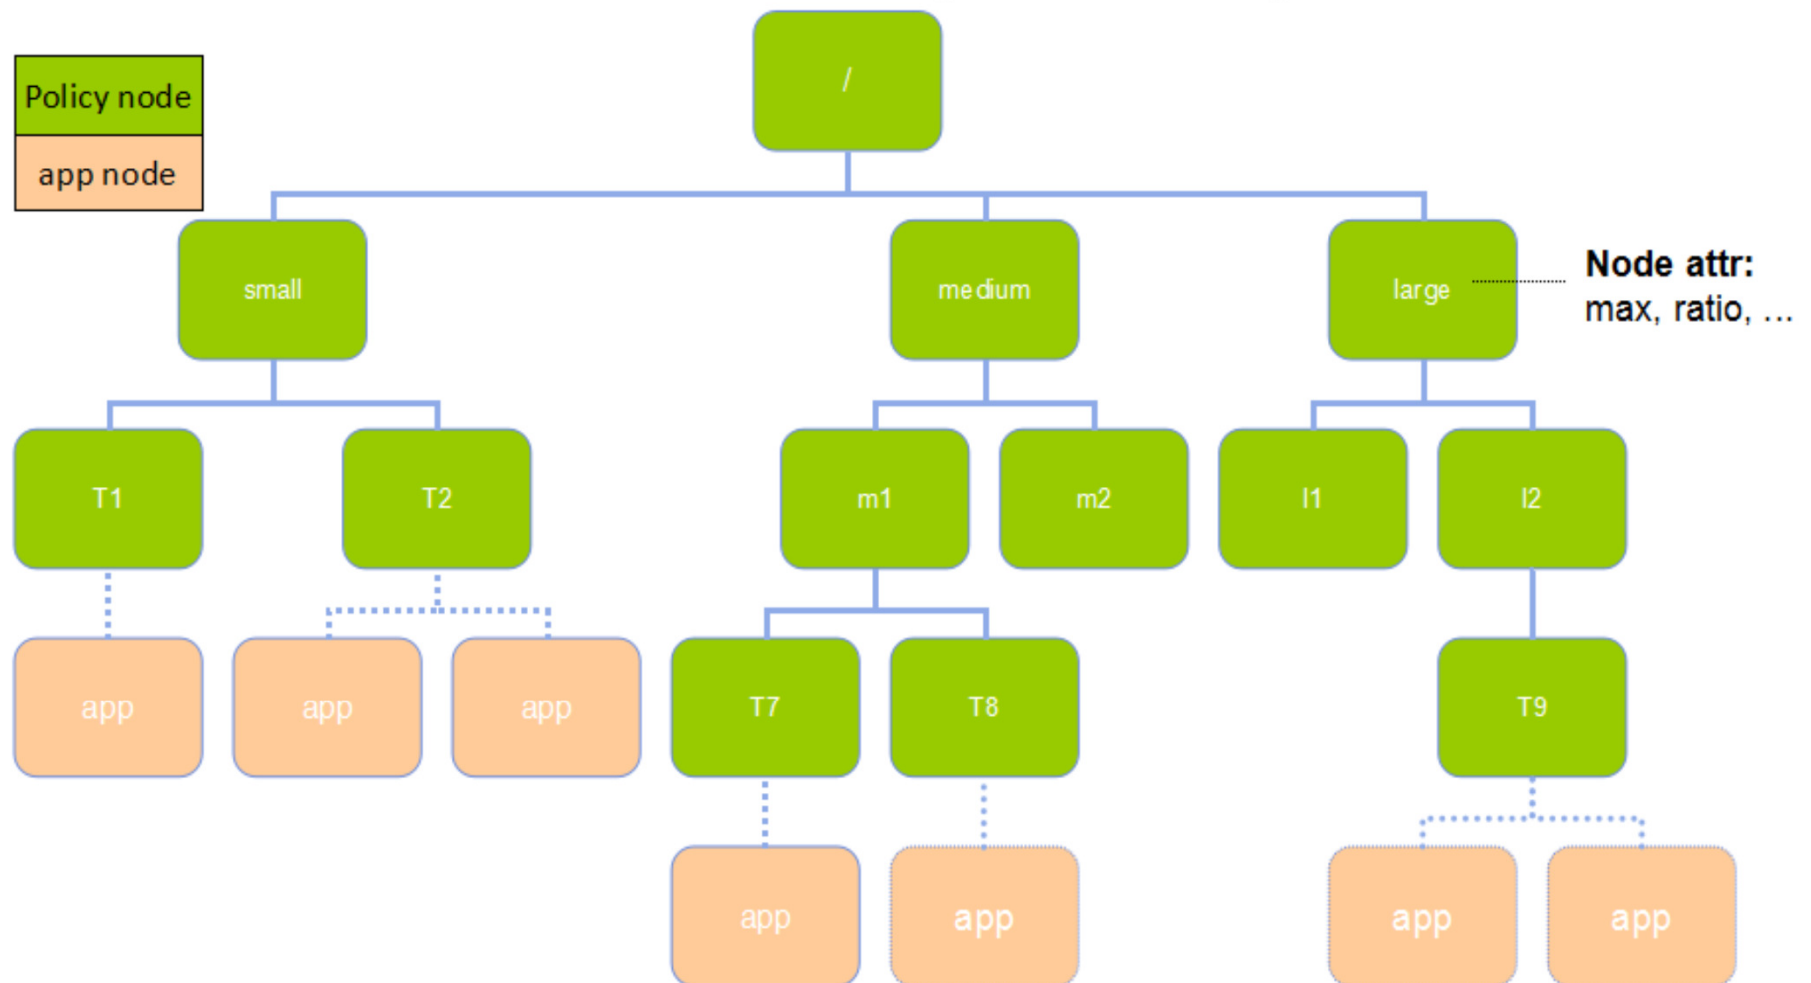




# Flexible scheduling policies (Spark service plans)



# Flexible hierarchical policy in Spark Session scheduler

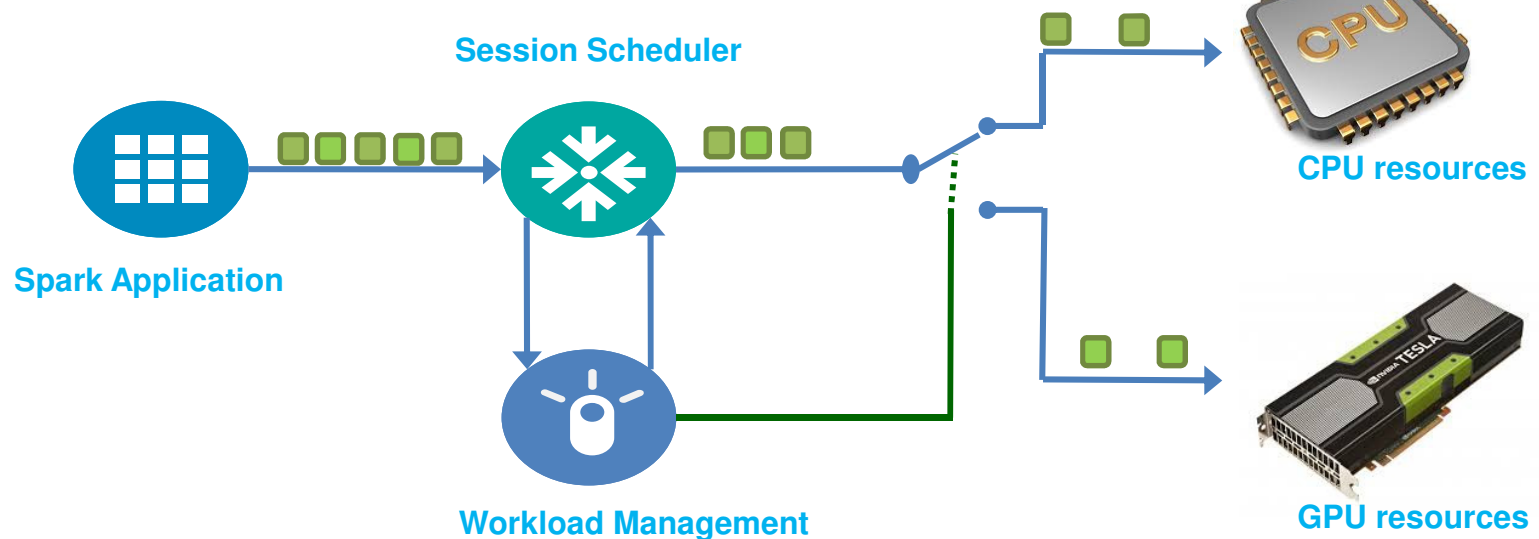


## Faster Time to Results – GPU Support

- Accelerate Spark applications with GPUs

– Presented at Spark Summit San Francisco June 2016

- Conductor scheduler interfaces with Spark scheduler to ensure that GPU resources are assigned to the applications that can use them



# Agenda

1

Bluemix Spark Cloud Technical Challenges

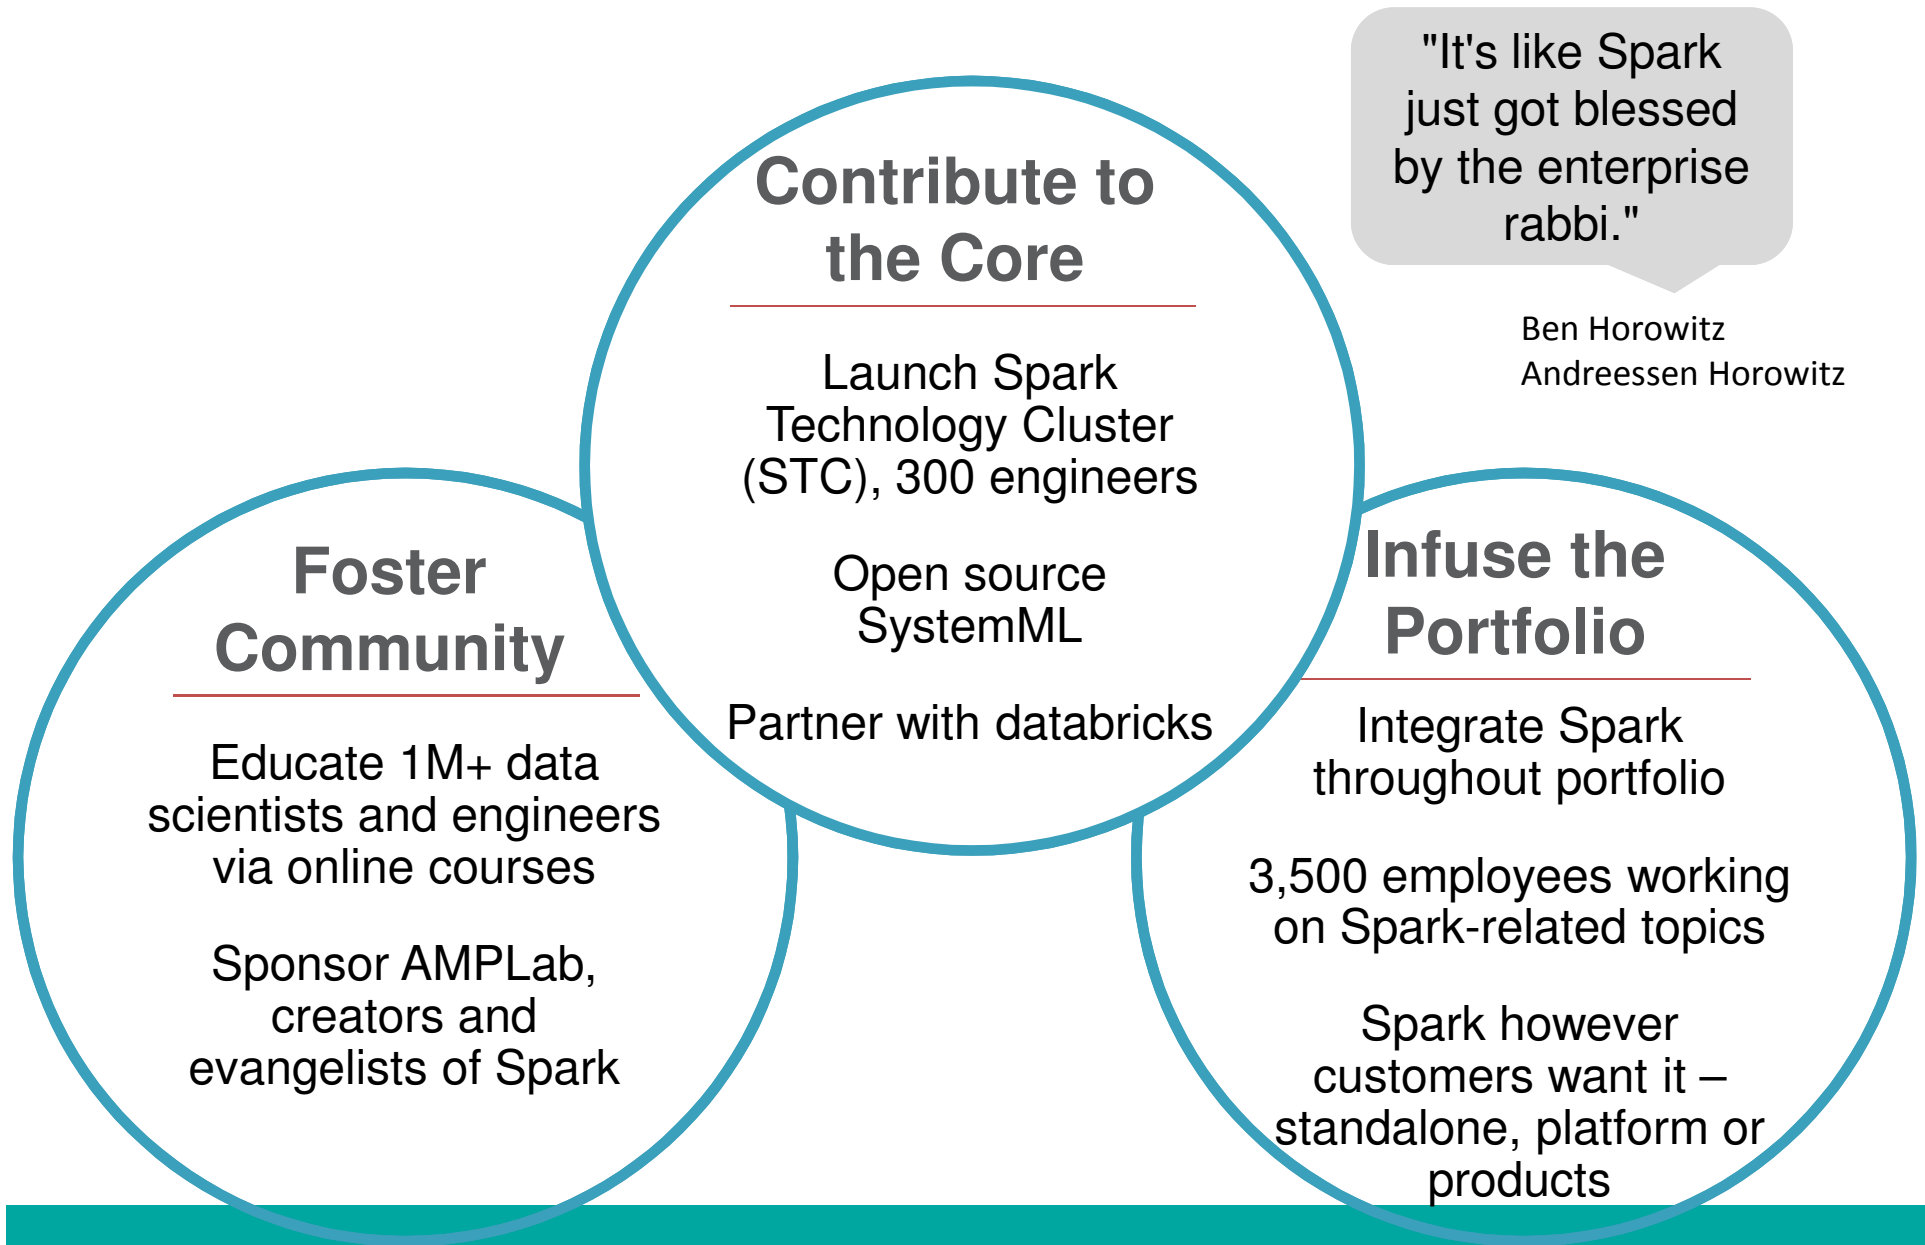
2

Bluemix Spark Cloud Architecture & Implementation

3

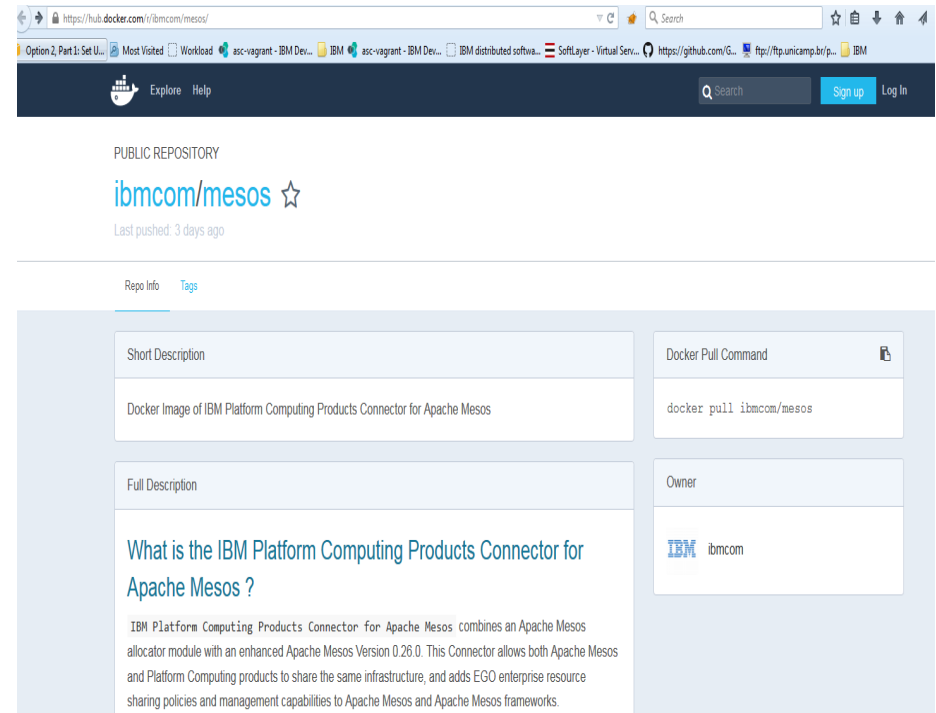
IBM – Spark & Mesos Community

# IBM is all-in on its commitment to Spark



# Mesos OSS Community Activities


- Active development with Mesos community – 11 IBM Developers.
- 100+ JIRAs delivered or in progress
- Leading or participating in several work streams: POWER Support, Optimistic Offers, Container Support, GPU Support, Swarm and Kubernetes integration
- Relationship with Mesosphere – weekly calls, on-site developer presence
- Attendance at MesosCon 2016 with sponsorship and booth
- ***Technical Preview of Mesos with IBM Value-Add on Docker Hub – Both x86 and POWER images***



## IBM Committed Mesos Patches


Note: The data is counted from Jun.2015 until now.

Commits per company



	Company	Commits ▼	LOC
1	Mesosphere	2,607	288,387
2	Twitter	295	35,031
3	Microsoft	235	20,899
4	IBM	224	19,388

Company Commit Trends



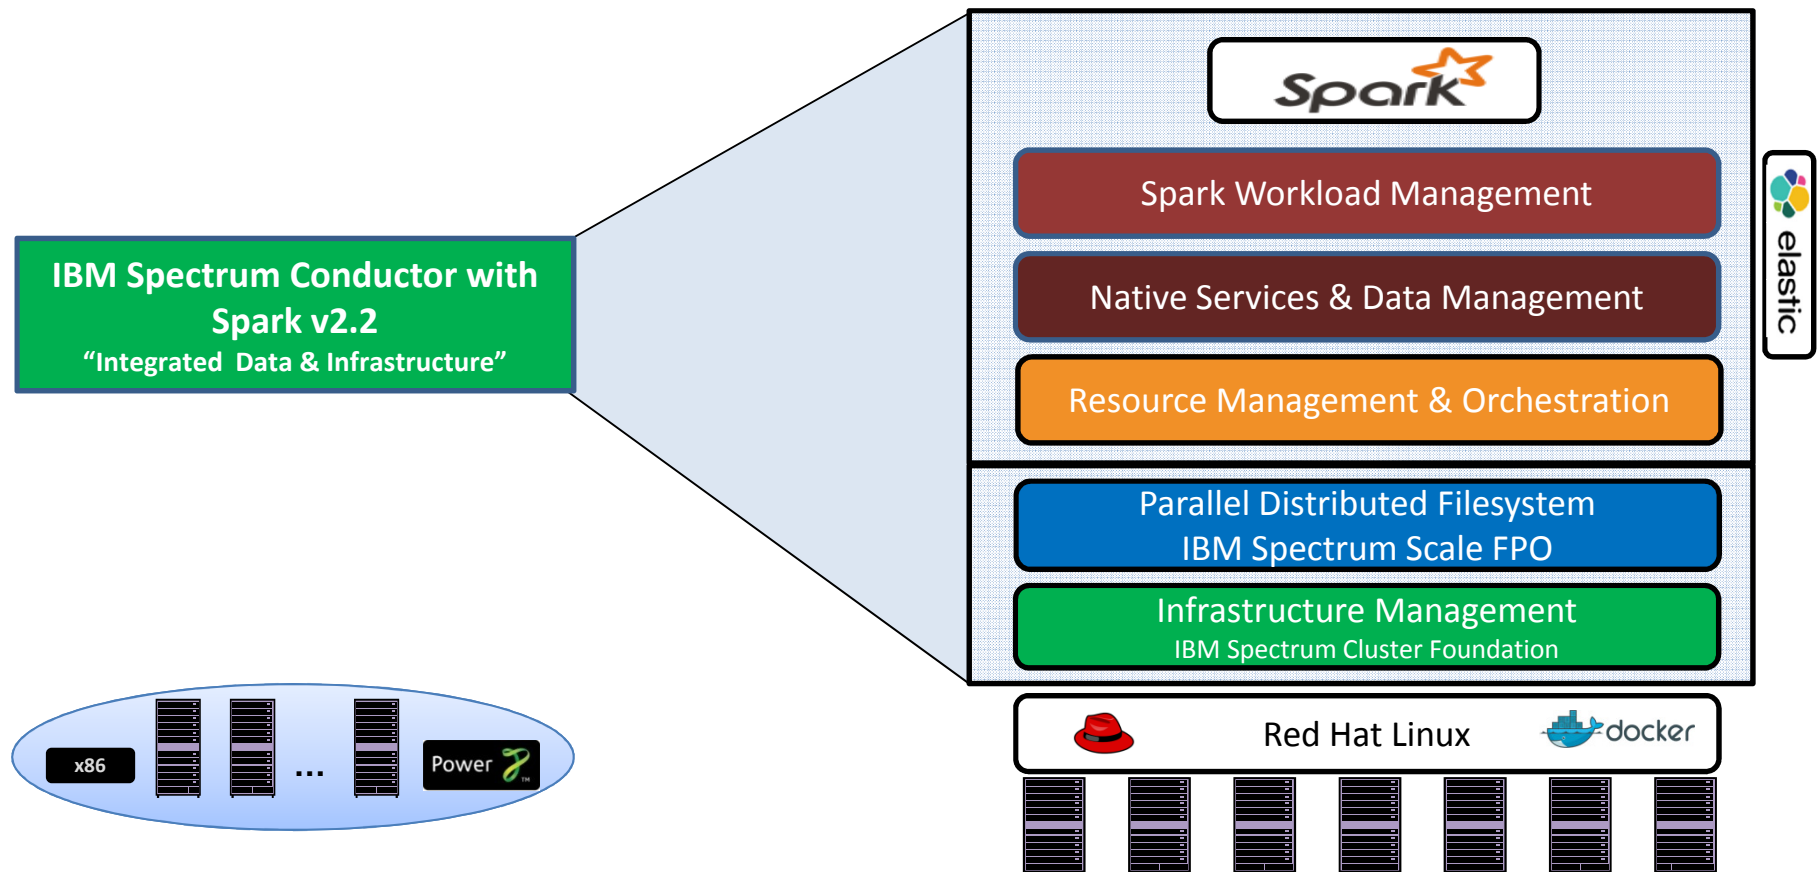
# Appendix

1. **IBM Analytics for Apache Spark**  
<https://console.ng.bluemix.net/catalog/services/apache-spark>
2. **IBM Data Science Experience**  
<http://datascience.ibm.com/>
3. **IBM DataWorks (Access, combine, transform – ACT)**  
<http://www.ibm.com/analytics/us/en/technology/cloud-data-services/dataworks/>



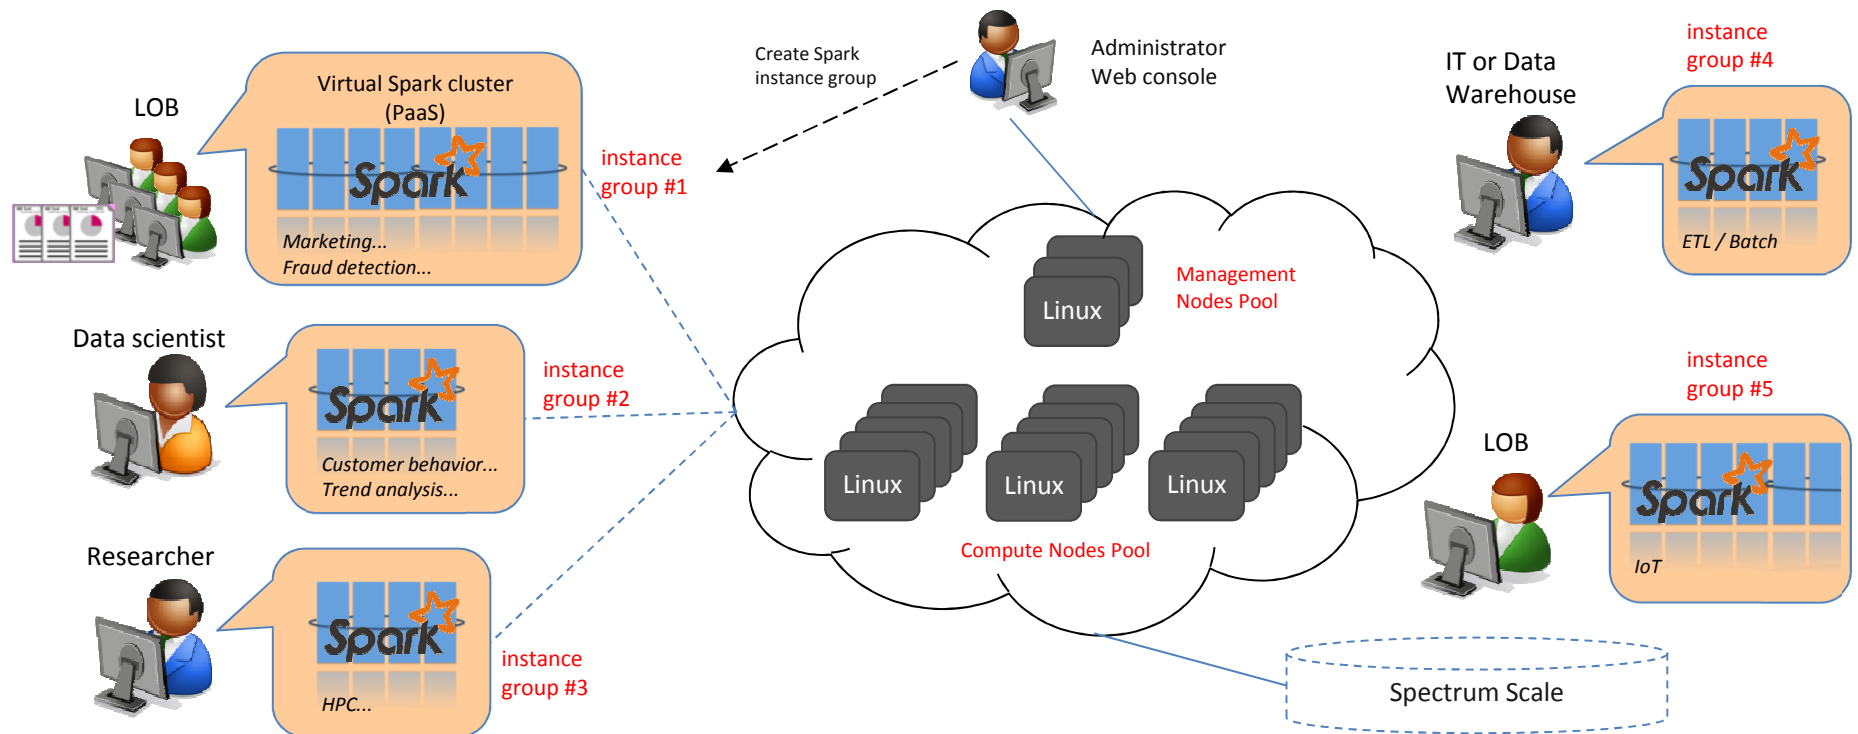
# Backup

## Apache Spark Cloud on-prem – Spectrum Conductor with Spark

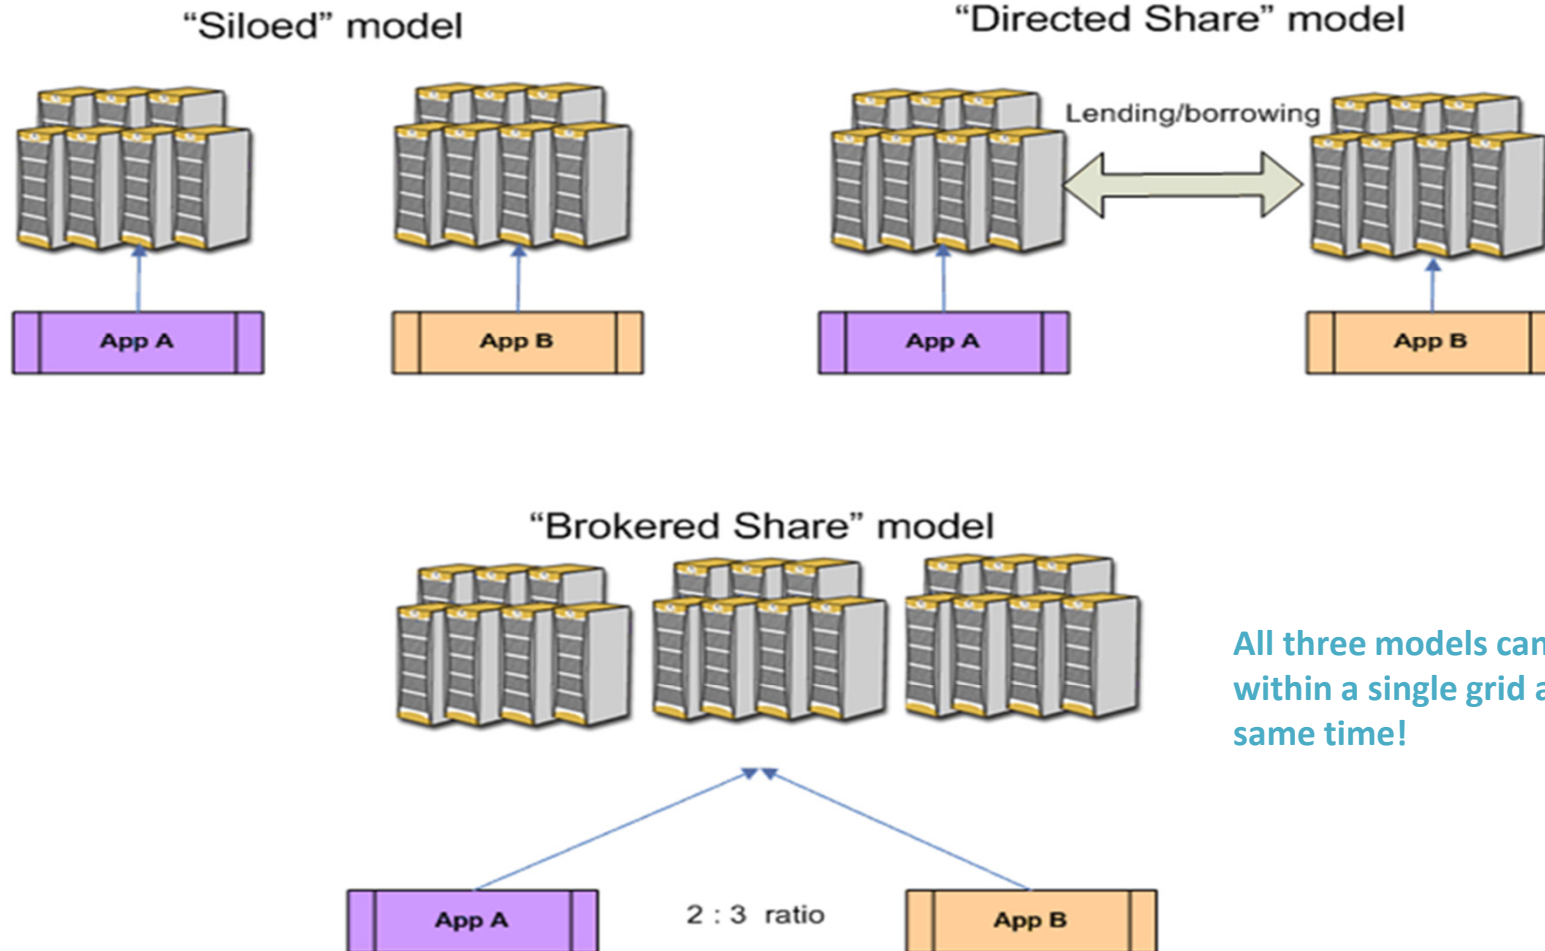


## Spark Shared Services Model – Private Spark Cloud

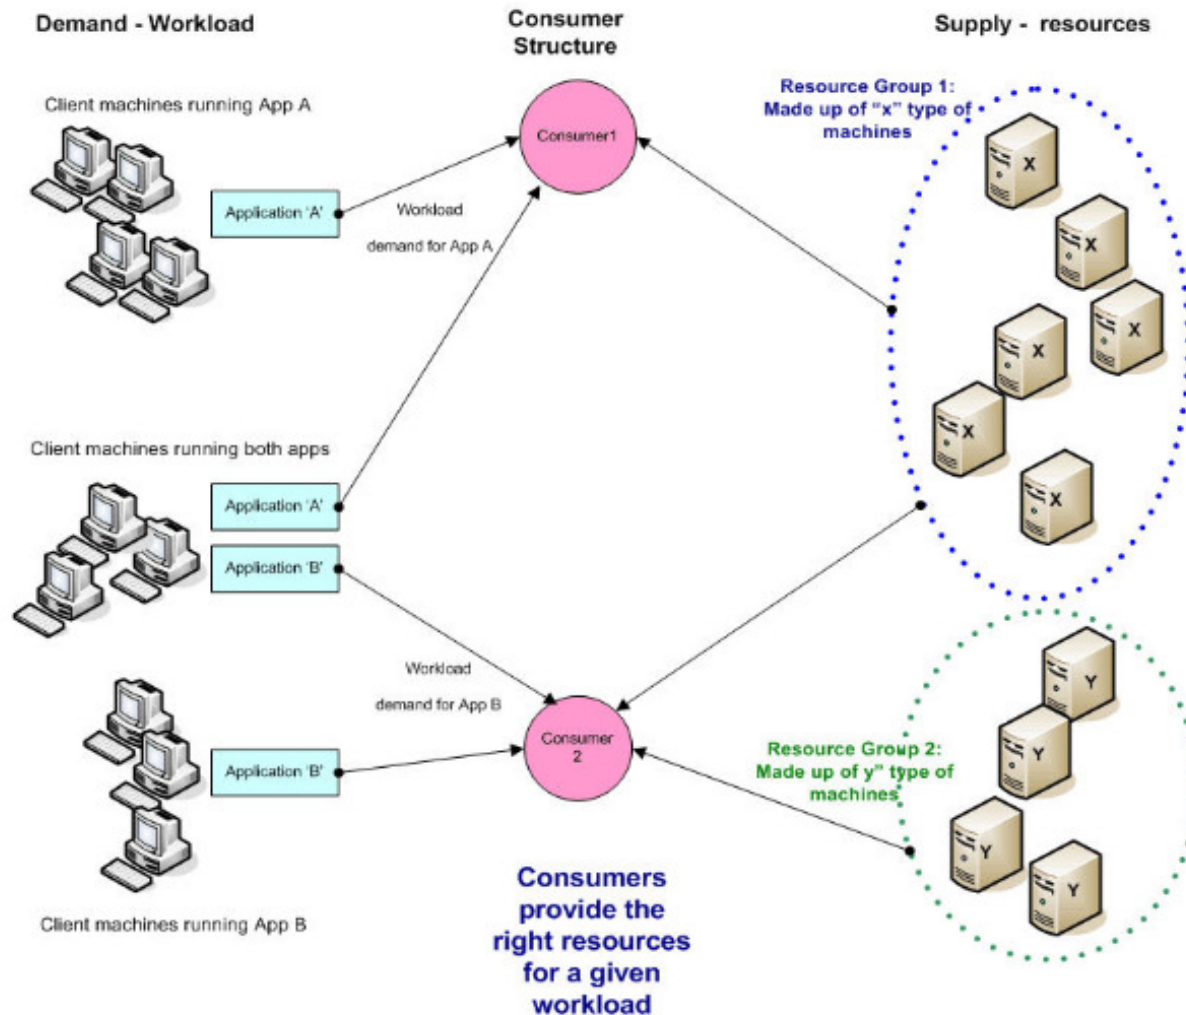
- Physical view: **Spectrum Conductor with Spark** installed on each Linux Server
- Logical view: Users (groups) have their own Spark cluster and they are isolated, protected, secured by Spark Instance Groups – Managed by SLA




## Resource Sharing Logic *Dynamic Runtime Sharing*



# Consumer Based Sharing Model



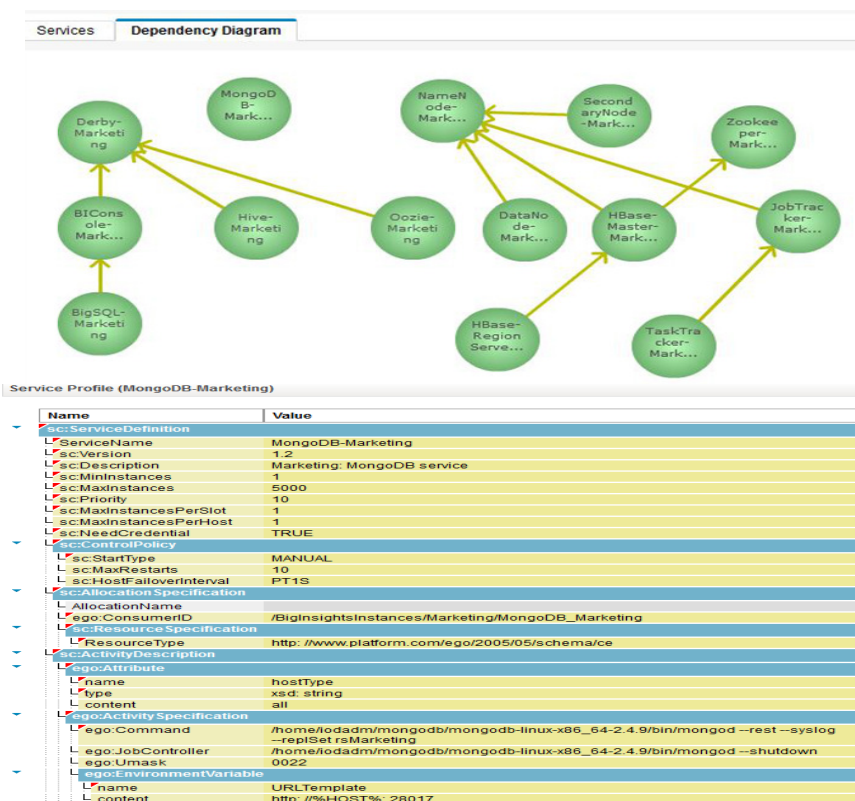
*Individual User, project, department, LOB, or entire company*



*Consumers can be divided into lower level consumers, which may be sub divided. The lowest level consumer is the level at which the application is associated.*

# IBM Spectrum Conductor – Complex Service Management

Support for complex application services



**A generalized service controller for complex long running application services**

- Service and application definition
- Service life cycle management
- Complex service dependency
- HA, Persistency, virtual IP mgmt
- Elastic service pool
- Multiple triggers for grow/shrink
- Dynamic services deployment
- Resource sharing among long running services and tasks/jobs
- Stateful vs. stateless services
- API & scriptable interface

# Spectrum Scale 4.2 Integration

- Multi-protocol support – Posix, HDFS, Object, etc.
- End-to-End, fully supported Spark solution with Spectrum Scale
  - No dependency to Hadoop / HDFS
- Benefit from all Spectrum Scale Enterprise data management features
  - Fileset based management
  - Information Lifecycle Management
  - Active file management
  - Retention and Immutability
- Integrated Spectrum Scale Monitoring

